

Reducing Code Size with Function Merging

Rodrigo Caetano de Oliveira Rocha



Doctor of Philosophy
Institute of Computing Systems Architecture
School of Informatics
University of Edinburgh
2021

Abstract

Resource-constrained devices for embedded systems are becoming increasingly important. In such systems, primary and secondary memories are highly restrictive, making code size in most cases even more important than performance. Compared to more traditional platforms, memory is a larger part of the cost and code occupies much of it. Despite that, compilers make little effort to reduce code size. One important optimisation for code-size reduction is function merging. This technique eliminates redundant code across functions by merging them into a single function.

However, production compilers only apply this optimisation to identical functions, while research compilers improve on that by merging the few functions with identical control-flow graphs and signatures. Overall, existing solutions are insufficient and we end up having to either increase cost by adding more memory or remove functionality from programs.

This thesis introduces the first techniques capable of merging arbitrary pair of functions. Our insight is that the weak results of prior function merging techniques are due to their rigid and overly restrictive solutions instead of the lack of duplicate code in the input program. Our solution corroborates this insight, resulting in drastic code size reductions while also reducing end-to-end compilation time.

First, we introduce FMSA, a novel technique that can merge arbitrary functions through sequence alignment, a bioinformatics algorithm for identifying regions of similarity between sequences. We combine this technique with an intelligent exploration mechanism to direct the search towards the most promising function pairs. Our approach is more than $2.4\times$ better than the previous state of the art, proposed by von Koch et al., reducing code size by up to 25%, with an overall average of 6%. FMSA increases end-to-end compilation time by an average of 15%.

While representing a leap forward, experiments show that FMSA fails to reduce code size in some cases where it would be intuitively expected to work. This limitation stems from its inability to directly handle phi-nodes. Instead, FMSA applies register demotion to replace all such nodes with memory operations, in an attempt to simplify the code generation processes. We build on this technique and develop SalSSA that fully supports the SSA form, removing any need for register demotion. By doing so, we notably increase the number of profitably merged functions. Experimental results on the SPEC 2006 and 2017 suites show that our approach delivers on average, 7.9% to 9.7% reduction on the final size of the compiled code. Moreover, as a result of aligning shorter sequences of instructions and reducing the number of wasteful merge

operations, our new approach incurs an average compile-time overhead of only 5%, $3\times$ less than FMSA, while also reducing memory usage by over $2\times$.

Finally, we continue to build on SalSSA by developing HyFM, which delivers similar levels of code size reduction for significantly lower compilation time and memory usage. To this end, we introduce an alignment strategy that works at the basic block level. Since basic blocks are usually much shorter than functions, even a quadratic alignment is acceptable. However, we also propose a linear algorithm for aligning blocks at a much lower cost. We extend this strategy with a multi-tier profitability analysis that bails out early from unprofitable merging attempts. By aligning individual pairs of blocks, we are able to decide their alignment’s profitability before actually generating code. Experimental results on SPEC 2006 and 2017 show that HyFM needs orders of magnitude less memory, using up to 48 MB or 5.6 MB, depending on the variant used, while SalSSA requires 32 GB in the worst case. HyFM also runs over $4.5\times$ faster, while still achieving comparable code size reduction. Combined with the speedup of later compilation stages due to the reduced number of functions, HyFM contributes to a reduced end-to-end compilation time.

Acknowledgements

TODO.

Declaration

I declare that this thesis was composed by myself, that the work contained herein is my own except where explicitly stated otherwise in the text, and that this work has not been submitted for any other degree or professional qualification except as specified. Some of the material used in this thesis has been published in the following papers:

- Rodrigo Rocha, Pavlos Petoumenos, Zheng Wang, Murray Cole, and Hugh Leather. “Function merging by sequence alignment.” In IEEE/ACM International Symposium on Code Generation and Optimization (CGO), pp. 149-163. Best Paper Award. 2019.
- Rodrigo Rocha, Pavlos Petoumenos, Zheng Wang, Murray Cole, and Hugh Leather. “Effective function merging in the SSA form.” In ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI), pp. 854-868. 2020.
- Rodrigo Rocha, Pavlos Petoumenos, Zheng Wang, Murray Cole, Hugh Leather, Kim Hazelwood. “HyFM: Function Merging for Free.” Submitted to the International Conference on Languages Compilers, Tools and Theory of Embedded Systems (LCTES). 2021.

(Rodrigo Caetano de Oliveira Rocha)

Table of Contents

1	Introduction	1
1.1	The Importance of Code Size for Different Domains	2
1.2	Limitations of Existing Function Merging	3
1.3	Contributions	4
1.3.1	Function Merging by Sequence Alignment	5
1.3.2	Effective Function Merging in the SSA Form	5
1.3.3	Function Merging for Free	6
1.4	Structure	7
2	Background	9
2.1	Compiler Infrastructure	9
2.1.1	Link-Time Optimisations	11
2.1.2	Target-Specific Cost Models	12
3	Related Work	13
3.1	Code-Size Optimisations	13
3.1.1	Constant Folding	13
3.1.2	Unreachable-Code Elimination	13
3.1.3	Dead-Code Elimination	14
3.2	Merging Identical Functions	14
3.2.1	Merging Identical Object Code During Link Time	16
3.2.2	Identical Function Merging	17
3.3	Merging Beyond Identical Functions	19
3.4	Code Factoring	22
3.5	Code Similarity	23
4	Function Merging by Sequence Alignment	25
4.1	Motivation	26

4.2	Our Approach	28
4.2.1	Overview	28
4.2.2	Linearization	30
4.2.3	Sequence Alignment	30
4.2.4	Equivalence Evaluation	31
4.2.5	Code Generation	32
4.3	Focusing on Profitable Functions	35
4.3.1	Profitability Cost Model	38
4.3.2	Link-Time Optimization	39
4.4	Evaluation	40
4.4.1	Experimental Setup	40
4.4.2	Code-Size Reduction	41
4.4.3	Compilation Overhead	45
4.4.4	Performance Impact	47
4.5	Conclusion	47
5	Effective Function Merging in the SSA Form	49
5.1	Motivating Example	50
5.2	The SalSSA Approach	54
5.2.1	Control-Flow Graph Generation	55
5.2.2	Operand Assignment	57
5.2.3	Preserving the Dominance Property	61
5.2.4	Phi-Node Coalescing	62
5.3	Evaluation	63
5.3.1	Experimental Setup	64
5.3.2	Evaluation on SPEC CPU	64
5.3.3	Evaluation on MiBench	66
5.3.4	Further Analysis	69
5.3.5	Memory Usage	69
5.3.6	Compilation Time Overhead	71
5.3.7	Performance Overhead	73
5.4	Conclusion	73
6	Function Merging for Free	75
6.1	Background and Motivation	76
6.1.1	Function Merging via Sequence Alignment	76

6.1.2	Limitations of SalSSA	77
6.1.3	When Less is More	79
6.2	Hybrid Function Merging	80
6.2.1	Overview	80
6.2.2	Pairing Similar Basic Blocks	81
6.2.3	Aligning Paired Basic Blocks	82
6.2.4	Multi-Tier Profitability Analysis	83
6.2.5	Independence from Code Layout	84
6.3	Evaluation	85
6.3.1	Experimental Setup	85
6.3.2	Code Size Reduction	86
6.3.3	Speeding Up Function Merging	88
6.3.4	End-to-End Compilation Time	91
6.3.5	Code Size and Compilation Time Trade-Off	91
6.3.6	Memory Usage	92
6.4	Conclusion	94
7	Conclusion	97
7.1	Contributions	97
7.1.1	Merging Arbitrary Pairs of Functions	97
7.1.2	Effective Code Generator for the SSA Form	98
7.1.3	Merging Functions One Block at a Time	99
7.2	Future Work	99
7.2.1	Handling Code Reordering at the Instruction Level	100
7.2.2	Merging Across All Scopes	101
7.2.3	Scaling for Large Programs	101
7.2.4	Powered by Deep Learning	101
7.2.5	Avoiding Performance Overheads	102
7.2.6	Less Memory Usage by JIT	102
	Bibliography	103

List of Figures

2.1	Overview of the three-phase compiler infrastructure.	9
2.2	Breakdown of the frontend, illustrating how compilers are organised as a series of phases.	10
2.3	Sequence of representations used during the compilation pipeline in modern compilers.	10
2.4	Overview of the default compilation pipeline.	11
2.5	Overview of the compilation pipeline using LTO.	12
3.1	Example of identical functions.	15
3.2	Two function extracted from the <code>447.dealIII</code> benchmark that are not identical at the source level, but after applying template specialisation and optimisations they become identical at the IR level.	15
3.3	Example of two pairs of highly similar functions. Because they are not identical, they cannot be merged by the function merging technique currently found in major compilers.	20
3.4	An example of two functions with isomorphic CFGs and their corresponding basic blocks arranged side by side. Instructions in paired basic blocks are compared in a pairwise manner.	21
3.5	An example of a merged basic block containing two mismatching pairs of instructions. A split is added for every pair of mismatching instructions with the phi-node instruction added to the their immediate point of convergence.	22
4.1	Example of two functions from the benchmark <code>sphinx</code> with different parameters that could be merged, as shown at the bottom. We highlight where they differ.	26

4.2	Example of two functions from the benchmark <code>libquantum</code> with different CFGs that could be merged, as shown at the bottom. We highlight where they differ.	27
4.3	Overview of our function-merging technique. Equivalent segments of code is represented in light green and the non-equivalent ones in dark red.	29
4.4	Linearizing the CFG of an example function.	30
4.5	The sequence alignment between two functions.	31
4.6	Example of a merge operation on the parameter lists of two functions.	33
4.7	Overview of our exploration framework.	36
4.8	Average CDF for the position of the profitable candidate and the percentage of merged operations covered. 89% of the merge operations happen with the topmost candidate.	37
4.9	In our experiments we use a compilation pipeline with a monolithic link-time optimization (LTO).	39
4.10	Object file size reduction for Intel (top) and ARM (bottom). We evaluate our approach (FMSA) under four different exploration thresholds, which control how many potential merging pairs we examine for each function before making a decision. Even for a threshold of one, we outperform the state-of-the-art by $2.4\times$ (Intel) and $1.9\times$ (ARM).	41
4.11	Object file size reduction for Intel on the Mibench benchmark suite. Our approach (FMSA) is the only one able to achieve a meaningful reduction on these benchmarks.	45
4.12	Compilation-time overhead on the Intel platform. For exhaustive exploration (not shown) the average overhead is $25\times$. Through ranking, we reduce overhead by orders of magnitude. For an exploration threshold of one, FMSA has an overhead of only 15%.	45
4.13	A compilation-time breakdown isolating the percentage for each major step of the optimization ($t=1$).	46
4.14	Runtime overhead on the Intel platform. Performance impact is almost always statistically insignificant. For the few benchmarks affected, FMSA merges hot functions.	46

5.1	Original input functions to be merged, before register demotion. These simplified functions highlight a problem commonly seen in real programs.	51
5.2	Desired merged function that can be produced by an expert. An extra argument called <code>%fid</code> is used to select between the two functions. This represents a gain of about 20% in the total number of instructions. . .	52
5.3	Aligned example functions after register demotion. The functions double in size after demotion, slowing down alignment. Merging some of the generated stack accesses will prevent eliminating them later through register promotion.	53
5.4	Average normalized function size, before and after register demotion, across all functions in each program from the SPEC 2006 benchmark suite. Register demotion increases function size by almost 75% on average.	54
5.5	Example functions aligned without register demotion. <i>Phi-nodes</i> are excluded from alignment.	55
5.6	Merged CFG produced by SalSSA. Code corresponding to a single input basic block may be transformed into a chain of blocks, separating matching and non-matching code. The generator inserts conditional and unconditional branches to maintain the same order of instructions from the input basic block. Operands and edges highlighted in blue will be resolved by the operand assignment described in Section 5.2.2.	56
5.7	Operand selection for the <code>call</code> instruction in L_4 from Figure 5.6. Mismatching operands chosen with a <code>select</code> instruction on the function identifier.	58
5.8	Optimizing operand assignment for commutative instructions. Example of a merged <code>add</code> instruction that can have its operands reordered to allow merging the two uses of <code>%m</code> , avoiding a <code>select</code> instruction. . . .	58
5.9	Label selection for mismatched terminator instruction operands L_{f1} and L_{f2} corresponding to labels of two different basic blocks. We handle control flow in a new basic block, L_{sel} with a conditional branch on the function identifier targeting the two labels. We use the label of the new block as the merged terminator operand.	59

5.10	Optimizing label assignment for conditional branches. Example of a merged <code>br</code> instruction that can have its label operands reordered, trading two label selections by one xor operation.	60
5.11	Landing blocks are added after operand assignment and are assigned to invoke instructions as operands.	60
5.12	Example of how SalSSA uses the standard SSA construction algorithm to guarantee the dominance property of the SSA form.	61
5.13	Phi-node coalescing reduces the number of phi-nodes and selections. .	62
5.14	Reducing the number of phi-nodes by coalescing disjoint definitions with no user instructions in common.	63
5.15	Compilation pipeline used for the evaluation. Both SalSSA and FMSA are applied in LTO mode.	64
5.16	Linked object size reduction over LLVM LTO when performing function merging with SalSSA or FMSA on SPEC CPU 2006 (a) and 2007 (b). Each approach was evaluated using three different exploration thresholds. On SPEC CPU2006, SalSSA reduces code size by 9.3% to 9.7% on average, almost twice as much as FMSA. On SPEC CPU2017, SalSSA reduces code size by 7.9% to 9.2% on average, more than twice as much as FMSA.	65
5.17	The percentual reduction in size of the linked object files, targeting the ARM architecture. We evaluate SalSSA or FMSA over the LLVM LTO on the MiBench embedded benchmark suite. Each approach was evaluated using three different exploration thresholds. SalSSA achieves a geo-mean reduction of 1.4% to 1.6%, about twice as much as FMSA.	68
5.18	A breakdown of SalSSA[$t = 1$] on the <code>djpeg</code> benchmark. The actual contribution to the final code size for each merge operation deemed profitable by the cost model.	68
5.19	Evaluation of the impact of phi-node coalescing on the size of the final object file. SalSSA-NoPC, which includes phi-node coalescing, has a measurable benefit over the alternative without phi-node coalescing (SalSSA-NoPC). When enabled, phi-node coalescing achieves up to 7% of code size reduction on top of SalSSA-NoPC.	69

5.20	Total number of profitable merge attempts for SalSSA and FMSA on 19 SPEC CPU2006 benchmarks. For both cases, we used the lowest exploration threshold ($t=1$). SalSSA achieves 31% more profitable merge operations.	70
5.21	Peak memory usage during compilation time on the SPEC CPU2006 benchmark. On average, SalSSA requires less than half the memory used by FMSA.	70
5.22	Speedup over the accumulated time spent on both sequence alignment and code generation. SalSSA produces significantly less overhead than the FMSA.	71
5.23	End-to-end compile-time for SalSSA and FMSA for three different exploration thresholds and 19 different SPEC CPU2006 benchmark. Compile-time is normalized to that of the baseline with no function merging. SalSSA reduces the overhead of function merging by $3\times$ to $3.7\times$ on average.	72
5.24	Comparison between the runtime impact from FMSA and SalSSA. Our approach increases the runtime overhead because it merges more functions. For most benchmarks, the overhead is small. For the rest, profiling-directed merging would eliminate the overhead.	72
6.1	Breakdown of the relative runtime for the different stages from SalSSA. Alignment takes 25 seconds and 4.2 minutes on <code>638.imagick_s</code> and <code>602.gcc_s</code> , respectively.	78
6.2	Example extracted from <code>483.xalancbmk</code> in SPEC CPU2006. Instructions marked green have been aligned through sequence alignment with an instruction from the other function. SalSSA would attempt merging all matched instructions but only the ones in fully aligned basic blocks would be profitable.	79
6.3	Two examples of the pairwise alignment. Only instructions in corresponding positions are aligned. Instructions match if they have the same opcode.	83
6.4	Example with code reordering extracted from the <code>450.soplex</code> program.	84
6.5	Linked object size reduction over LLVM LTO when performing function merging with HyFM or SalSSA on SPEC CPU 2006 and 2017. On average, HyFM improves code size reduction.	87

6.6	Speedup of the function merging pass in isolation relative to SalSSA. The multi-tier profitability analysis reduces the number of unprofitable merge operations leading to a significant speedup.	89
6.7	Breakdown of the relative runtime for the different stages of the function merging pass. All measurements are normlized by SalSSA’s total runtime on the corresponding benchmark. For every benchmark, we show SalSSA, [PA,NMP], [PA], [NW,NMP], and [NW], in this order.	90
6.8	Normalized end-to-end compilation time for SPEC 2017 and SPEC 2006 relative to LLVM LTO.	92
6.9	Average reduction speed on both SPEC 2006 and 2017.	93
6.10	Peak memory usage of SalSSA and HyFM variants for SPEC 2006 and 2017 in log scale. SalSSA has a peak memory usage several orders of magnitude hundreds higher than all other approaches. The pairwise alignment variants of HyFM need on average only a seventh of the memory needed by the Needleman-Wunsch variants.	95
7.1	Example of how even trivial reordering is poorly handled by the existing solutions.	100

List of Tables

4.1	Number and size of functions present in each SPEC CPU2006 benchmark just before function merging, as well as number of merge operations applied by each technique.	42
4.2	Number and size of functions present in each MiBench benchmark just before function merging, as well as number of merge operations applied by each technique.	44
5.1	Number and size of functions present in each MiBench benchmark just before function merging, as well as number of merge operations applied by each technique.	67

Chapter 1

Introduction

While often overlooked, program size can be a first-order constraint. From tiny embedded devices up to cloud servers, these systems are all operating under limited addressable memory, storage, or bandwidth. When the program becomes excessively large relative to the given constraints, this has a detrimental effect on the system. This is very likely to happen as programs gain new features over time, continuously growing in size and complexity [8, 38]. In such scenarios, reducing the application footprint is essential [4, 8, 31, 52, 53, 61].

Despite the importance of keeping code size small, compilers still make little effort to reduce it, except for some classical optimisations, such as dead-code elimination. Their efforts are usually limited to disabling performance optimisations that tend to increase size, such as loop unrolling or inlining. Developers might have more luck just removing functionality from their libraries [31] or hand-optimizing their code [62]. However, these efforts are often undesirable if even possible at all. Therefore, we must develop and tune compilation techniques primarily focused on reducing code size.

Code-size optimisations work by replacing a piece of code with another that is semantically equivalent but uses fewer or smaller instructions, in the binary format, sometimes combining and reusing equivalent pieces of code. Classical optimisations that are effective in reducing code size include the elimination of redundant, unreachable, or dead code [6, 9, 15]. Although initially motivated by performance, these classical optimisations achieve better performance by reducing the static number of instructions in the code, which translates to fewer dynamic instructions during runtime.

Recently, we have seen some progress with optimisations based on merging equivalent code within or across functions [8, 19]. One important optimisation capable of reducing code size is function merging. In its simplest form, function merging reduces

replicated code by combining multiple identical functions into a single one [2, 39]. This optimisation is found in linkers, by the name of *identical code folding* (ICF), where text-identical functions at the bit level are merged [1, 35, 56].

It is obvious how function merging can reduce code size by removing function duplicates. Nevertheless, it can also potentially reduce compilation time. For example, when merging two identical functions, the remaining compilation pipeline will have one fewer function to process and optimise. These benefits are not as obvious when merging non-identical functions as it introduces extra code, adding complexity to the merged function. In the thesis, even though our main goal is reducing code size, we will focus on both dimensions: code size and compilation time.

1.1 The Importance of Code Size for Different Domains

In this section, we discuss in detail the importance of code size for different domains.

The embedded system market is rapidly growing. Embedded systems need to perform increasingly complex tasks, with their application binaries often reaching several megabytes in size, while running on inexpensive and resource-constrained devices. As a result, permanent storage and memory size becomes a limiting factor [45]. Just adding more memory is not always a viable option. Highly integrated systems-on-chip are common in this market and their memories typically occupy the largest fraction of the chip area, contributing to most of the overall cost. Even small increases in memory area translate directly to equivalent increases in cost, which lead to enormous levels of lost profit at large scales [18]. In addition to cost, embedded systems are also often limited by other factors such as weight, area, and energy consumption [20, 57]. All these factors limit the size of the storage and memory available in embedded systems.

Modern mobile applications tend to have large binaries that need to support as many devices as possible, including low-end devices with limited resources [21, 25]. Furthermore, Apple App Store imposes a limit when downloading an application over the mobile broadband. Applications larger than this limit must be downloaded only over the Wi-Fi. These restrictions on the size of the application may significantly impact revenues for critical businesses. Chabbi et al. [8] have shown that certain large applications may have over 90% of its total size being taken by their binary code, where the remaining size is due to media and resources. For these applications, reducing the application's binary size becomes of utmost importance for their businesses.

Beyond just mobile and embedded systems, powerful machines can also be unable

to properly handle extremely large programs. For example, compilation and load time can become impractical for extremely large programs and codebases [24, 29]. Moreover, address space also limits how large programs can become. Upgrading computers at scale is a challenging and costly process even for large datacenters [43, 63]. As a result, outdated 32-bit machines have an addressable memory space limited to less than 4 GB, setting a limit on program size. This limitation is even worse on machines with shorter word widths. In such constrained scenarios, reducing the application’s footprint is essential [4, 31, 52, 53, 61].

1.2 Limitations of Existing Function Merging

Google developed an optimization for the *gold* linker that merges functions that are identical at the bit-level [35, 56]. Similar machine-level implementations are also offered by other production compilers and linkers, such as MSVC [1]. However, such solutions are platform-specific and need to be adapted for each object code format and hardware architecture. Alternatively, compilers also provide a similar optimisation for merging identical functions at their mid-level intermediate representation (IR) which is therefore agnostic to the target hardware [2, 39]. Unfortunately, these optimisations can only merge fully identical functions with at most type mismatches that can be losslessly cast to the same format. These techniques can leverage their simplicity to efficiently identify groups of mergeable functions. First they compute the hash of all functions, then a tree structure is used to group equivalent functions based on their hash values.

More advanced approaches can identify similar, but not necessarily identical, functions and replace them with a single function that combines the functionality of the original functions while eliminating redundant code. At a high level, the way this works is that code specific to only one input function is added to the merged function but made conditional to a function identifier, while code found in both input functions is added only once and executed regardless of the function identifier. The function-merging technique presented by von Koch et al. [19] exploits similarity among functions. Two functions are structurally similar if both their function types are equivalent and their control-flow graphs (CFGs) are isomorphic. Two function types are equivalent if they agree in the number, order, and types of their parameters as well as their return types, linkage type, and other compiler-specific properties. In addition to the structural similarity of the functions, their technique also requires that corresponding

basic blocks have exactly the same number of instructions and that corresponding instructions must have equivalent resulting types. Mergeable functions are only allowed to differ in corresponding instructions, where they can differ in their opcodes or input operands.

Unfortunately, existing approaches fail to produce any noticeable code size reduction. In this work, we introduce a novel way to merge functions that overcomes major limitations of existing techniques. Our insight is that the weak results of existing function merging implementations are not due to the lack of duplicate code but due to the overly restrictive algorithms they use to find duplicates.

Our approach is based upon the concept of sequence alignment, developed in bioinformatics for identifying functional or evolutionary relationships between different DNA or RNA sequences. Similarly, we use sequence alignment to find areas of functional similarity in arbitrary function pairs. Aligned segments with equivalent code are merged. The remaining segments where the two functions differ are added to the new function too but have their code guarded by a function identifier. This approach leads to significant code size reduction.

Attempting to merge all pairs of functions is prohibitively expensive even for medium sized programs, considering the quadratic nature of sequence alignment. To counter this, our technique is integrated with a ranking-based exploration mechanism that efficiently focuses the search to the most promising pairs of functions. As a result, we achieve our code size savings while introducing little compilation-time overhead.

Compared to identical function merging, we introduce extra code to be executed, namely the code that chooses between dissimilar sequences in merged functions. A naive implementation could easily hurt performance, e.g by merging two hot functions with only a few similarities. We show that it is also possible to avoid performance degradation by incorporating profiling information in the decision-making, enabling the compiler to avoid merging functions that contain hot code.

1.3 Contributions

In this section, we provide an overview of the main contributions presented in this thesis. We provide the first techniques capable of merging arbitrary pair of functions. Then, we build on this technique to achieve an even greater code size reduction while also reducing its compilation-time overheads. Our final contribution is a function merging technique that can effectively reduce code size as well as end-to-end

compilation time.

1.3.1 Function Merging by Sequence Alignment

Existing function merging techniques fail to produce any noticeable code size reduction in most programs. Our insight is that the weak results of existing function merging implementations are not due to the lack of duplicate code but due to the rigid, overly restrictive algorithms they use to find duplicates. These techniques only work on either identical or mostly identical functions, limiting their potential to reduce code size.

As the first contribution of this thesis, we introduce a novel way to merge functions that lifts most of the restrictions imposed by prior techniques. Our technique is the first that allows merging arbitrary functions, including functions with different signatures and control flow graphs. The proposed optimisation uses sequence alignment to identify code similarity and guide the merging operation. It also merges parameters based on their usage, minimising the number of parameters and operand selection, and handles different return types using a union-like approach. The goal is to maximise the amount of merged code while minimising the overhead required to handle the differences.

For functions with little to no similarity, merging them might increase code size. However, because our technique is able to merge any pair of functions, it is necessary to identify which pairs of functions are the most profitable to merge. To this end, we introduce a novel ranking mechanism for focusing our optimisation on function pairs that are more likely to be profitably merged. The proposed mechanism first pre-computes fingerprints summarising each function and later use them to rank function candidates based on the fingerprint similarity. For each function, merging will be only attempted for the top ranked candidates, significantly reducing compilation overhead while still resulting in a meaningful code size reduction.

These contributions have been previously described in the following publication:

- Rodrigo Rocha, Pavlos Petoumenos, Zheng Wang, Murray Cole, and Hugh Leather. “Function merging by sequence alignment.” In IEEE/ACM International Symposium on Code Generation and Optimization (CGO), pp. 149-163. 2019.

1.3.2 Effective Function Merging in the SSA Form

Although our technique, introduced in Chapter 4, achieves impressive results, it does not directly handle phi-nodes which are fundamental to the SSA form. Instead, it

applies register demotion to replace all such nodes with memory operations, in an attempt to simplify the code generation processes. As we show in Chapter 5, after register demotion, functions tend to be almost twice as long due to an excessive number of memory operations. Therefore, such a strategy comes at the cost of poor merge results, larger memory footprint, and longer compilation time.

Chapter 5 presents our improved technique that avoids this pitfall with a new code generator capable of handling phi-nodes properly and completely bypassing register demotion. This approach results in better merged functions by not relying on later reversing the effects of register demotion. Merging smaller functions also significantly reduces compilation overhead, due to the quadratic nature of sequence alignment.

The code generator also includes a novel phi-node coalescing optimisation tailored for function merging, reducing the total number of phi-nodes and easing the pressure on registers. Phi-node coalescing is able to reduce the number of phi-nodes and selections, producing smaller merged functions and reducing code size even further.

These contributions have been previously described in the following publication:

- Rodrigo Rocha, Pavlos Petoumenos, Zheng Wang, Murray Cole, and Hugh Leather. “Effective function merging in the SSA form.” In ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI), pp. 854-868. 2020.

1.3.3 Function Merging for Free

Our solution presented in Chapter 5, SalSSA, achieves on average a 10% code size reduction but at the cost of crippling compile-time inefficiencies, especially for large programs. In Chapter 6, we show that SalSSA can lead to 40% slower compilation, taking up to 32 GB of peak memory usage even for modestly-sized program. Such a resource requirement is beyond what is typically available to a developer and thus unsuitable for optimizing real-life programs. These inefficiencies stem directly from its quadratic sequence alignment used to identify mergeable instructions in a pair of input functions.

In order to address these inefficiencies, we propose a new sequence alignment strategy that works on the basic block level. Because basic blocks tend to be much smaller than whole functions, this strategy greatly reduces the impact of the quadratic sequence alignment algorithm. For even greater speedups, we propose a linear pairwise alignment that works on pairs of basic blocks of the same size. This technique is often

enough to achieve good code size reduction, because profitably merged functions tend to have highly similar basic blocks.

We also propose a multi-tier profitability analysis. This include a fine-grain analysis that estimates the profitability of the aligned basic blocks before actually generating their merged code, allowing the compiler to bail out early from unprofitable merging attempts, speeding up the optimization process. This fine-grain analysis on the aligned blocks result in improved code size reduction.

We show that this technique is capable of reducing end-to-end compilation time. This result can be achieved due to two main reasons: 1) By merging functions and removing code duplicates, it reduces the overall amount of code that needs to be optimised and processed by the remaining compilation pipeline. 2) The compilation overhead required for merging functions is smaller than its benefits, resulting in a overall reduction on the end-to-end compilation time.

These contributions have been previously described in the following paper, which is currently under review:

- Rodrigo Rocha, Pavlos Petoumenos, Zheng Wang, Murray Cole, Hugh Leather, Kim Hazelwood. “HyFM: Function Merging for Free.” Submitted to the International Conference on Languages Compilers, Tools and Theory of Embedded Systems (LCTES). 2021.

1.4 Structure

This thesis is organised as follows:

Chapter 2 provides the main background. It provides an overview of compiler architecture and compiler optimisations for code size reduction. It also provides terminology and describes the sequence alignment algorithm from bioinformatics. Finally, it describes the deep learning techniques used in this work.

Chapter 3 surveys the relevant literature. First we describe the classic compiler optimisations that work by reducing code size. Then, we describe in detail the existing techniques for function merging. Finally, we discuss other techniques related to identifying code similarity.

Chapter 4 describes our novel function merging technique based on sequence alignment. It also presents its accompanying search strategy, where a fingerprint-

based ranking mechanism is used to focus the optimisation on functions with more similarities.

Chapter 5 describes our new code generator for function merging that is capable of effectively handling functions in the SSA form.

Chapter 6 describes our new alignment strategies that work on the basic block level, reducing compilation time overheads. This enables a fine-grain multi-tier profitability analysis capable of bailing out early from unprofitable merging attempts.

Chapter 7 summarises the overall findings of the thesis and outlines potential avenues for future research.

Chapter 2

Background

2.1 Compiler Infrastructure

Compilers are programming tools responsible for translating programs in a given source language to a lower-level target language. This compilation process must preserve the program semantics. Moreover, compilers are also expected to produce a good quality representation of the program in the target language, optimising for a given objective function. An important objective function is code size, i.e., the optimisation goal is to produce a representation of the program as small as possible.

Compilers are usually described as *three-phases* architecture, as shown in Figure 2.1: frontend, optimiser, and backend. The frontend is responsible for parsing, validating and diagnosing errors in the source code. This parsed source code is then translated into an intermediate representation, namely, the LLVM IR [36]. The optimiser is responsible for doing a broad variety of transformations, that are usually independent of language and target machine, to improve the code's performance. The backend, also known as the code generator, then translates the code from the intermediate representation onto the target instruction set. It is common for the backend to also perform some low-level optimisations that take advantage of unusual features of the supported architecture.




Figure 2.1: Overview of the three-phase compiler infrastructure.

However, these *three-phases* represent only a simplified view of their designed. In order to manage the complexity involved in optimising compilers, modern compilers

are actually designed in a highly modular manner, where they are organised as a series of phases that sequentially analyse and transform the program being compiled. For example, the frontend alone is subdivided into multiple phases, as shown in Figure 2.2. The *lexer* is responsible for tokenising the input stream of characters from the source code. This token stream is then consumed by the *parser*, producing a language-specific *Abstract Syntax Tree* (AST). The AST is an intermediate representation used during the semantic analysis, before being lowered to another intermediate representation.






Figure 2.2: Breakdown of the frontend, illustrating how compilers are organised as a series of phases.



(a) An overview of representations and their level of abstractions used during the compilation pipeline.



(b) An example of the sequence of representations used in real compilers for different programming languages.

Figure 2.3: Sequence of representations used during the compilation pipeline in modern compilers.

Moreover, several intermediate representations, with different levels of abstraction, are used throughout the compilation process from the source to the target language. Figure 2.3a illustrates the sequence of representations used by modern compilers, each one having a progressively lower level than the previous one [37]. The source language is parsed into an AST, which is then commonly lowered to another high-level intermediate representation (HIR). As shown in Figure 2.3b, this HIR is usually a language-specific IR, such as the Swift Intermediate Language (SIL), which is used

to solve domain-specific problems. A popular mid-level intermediate representation (MIR) is the LLVM IR, which is shared among many compilers. In the LLVM compiler infrastructure, the LLVM IR is later lowered into a low-level representation (LIR), called the Machine IR, which is finally lowered into the target assembly language. The machine-code generation alone might actually involve multiple other intermediate representations depending on the backend being used.

Different analyses and optimisations are better modelled at different abstraction levels. However, most analyses and optimisations are language- and target-independent and therefore are applied on the mid-level IR, maximising their reuse. An optimisation developed for the mid-level IR can be used with any supported programming language or target assembly.

2.1.1 Link-Time Optimisations

Compilers normally operate on a single translation unit at a time, where a translation unit in C or C++ represents a single source file and its expanded headers. Each translation unit is optimised separately and compiled into a single native object file. Finally, the linker combines multiple object files into a resulting binary or library. Figure 2.4 illustrates the standard pipeline for the compilation of multiple source files.




Figure 2.4: Overview of the default compilation pipeline.

However, this approach limits the impact of inter-procedural optimisations (IPO) to within each individual translation unit. For example, inlining can only occur between two functions in the same translation unit. In order to fully benefit from IPO, these optimisations should operate on the whole programs, instead of a single translation unit. When the optimising across all translation units being linked into an executable, the compiler can perform more aggressive optimisations that rely on whole-program information [30].

Figure 2.5 illustrates a common mechanism for enabling whole-program optimisation called link-time optimisation (LTO). In this approach, optimisations are applied




Figure 2.5: Overview of the compilation pipeline using LTO.

in two different moments. First, we have *early* optimisations being applied on a per translation unit basis. However, we also have *late* optimisations being applied after all translation units are linked together, allowing important inter-procedural optimisations to be applied on the whole-program.

In the LTO mode, compilers usually delay the generation of native object files until after the late optimisations. As shown in Figure 2.5, all translation units are linked while still in an IR better suited for optimisation.

2.1.2 Target-Specific Cost Models

Compilers have cost models to compare two pieces of code in order to decide whether or not a given code transformation is profitable. Therefore, these cost models play an essential role in the decision making of most compiler optimisations [46, 60].

These cost models provide target-dependent cost estimations, with respect to a given metric, approximating the cost of an IR instruction when lowered to machine instructions. Compilers usually provide multiple cost models, including: a performance cost model that estimates the latency of each instruction; a code-size cost model that estimates the binary size of each instruction. However, these cost models need not predict precisely the latency or binary size of each instruction, because, ultimately, these costs are used only for comparing two pieces of code in order to guide optimisations. Therefore, the cost estimated for one instruction is only relevant relative to other instructions.

Chapter 3

Related Work

3.1 Code-Size Optimisations

Although initially motivated by performance, many of the classical optimisations achieve better performance by reducing code size. A small code, besides having fewer instructions to execute, can also have a positive impact on the cache utilisation. Classical optimisations that are effective in reducing code size include the elimination of redundant, unreachable, and dead code, as well as certain kinds of strength reduction [6, 9, 15]. In this section, we will describe some of these classical size-reducing optimisations.

3.1.1 Constant Folding

Constant folding is an optimisation that operates on the instruction level, identifying instructions whose operands are constant values, performing the evaluation of the instruction at compile time, and replacing it by the resulting value. The effectiveness of constant folding can be augmented by combining it with constant propagation. Constant folding reduces code size by eliminating instructions that can be computed at compile time. Moreover, constant folding also works as an enabler to other optimisations, such as unreachable-code elimination (Section 3.1.2).

3.1.2 Unreachable-Code Elimination

Some functions may contain code that is unreachable. A code is unreachable if there is no valid control-flow path from the function's entry point that leads to it. Since unreachable code is guaranteed to never be executed, compilers should remove it to avoid code bloat.

Often, unreachable code is uncovered by other optimisations. For example, after constant propagation and constant folding, a conditional branch could have its condition evaluating to a constant, eliminating a path to one of its successor basic blocks. If no other path leads to that basic block, it becomes unreachable.

The algorithm to eliminate unreachable code works in a mark-sweep manner, performing two passes over the basic blocks of the CFG. The reachability analysis optimistically assumes that all basic blocks are dead until proven otherwise. First, it marks all blocks as unreachable. Next, starting from the entry point, it marks each block that it can reach as reachable. If all branches and jumps are unambiguous, then all unmarked blocks can be deleted. With ambiguous branches or jumps, the compiler must preserve any block that the branch or jump can reach. This analysis is simple and inexpensive.

3.1.3 Dead-Code Elimination

A value definition is dead if it is not used on any path from the point in which it is defined to the exit point of the function. In a similar way, an instruction is dead if it computes only values that are not used on any execution path leading from the instruction. Any dead definition or instruction can be simply removed without altering the program's semantics, therefore reducing code size [42].

The algorithm to eliminate dead code has some similarities with that for unreachable-code elimination described in Section 3.1.2. This algorithm also works in a mark-sweep manner [59]. First, the algorithm marks *critical* instructions as *alive*. An instruction is *critical* if it has an observable effect, for example, if it is a return instruction, a branch instruction, a function call, a memory operation, any instruction with side effect, etc. Then, the algorithm follows the *use-def chain* of every alive instruction, marking the operand instructions as alive. This process continues until no more instructions can be marked as alive. Finally, the sweep phase removes all instructions that have not been marked as alive, reducing code size.

3.2 Merging Identical Functions

In this section, we will discuss existing optimisations for merging identical functions. Figure 3.1 illustrates how identical functions can appear in real programs. The first pair of functions, shown in Figure 3.1a, were extracted from the `482.sphinx3` benchmark.

The only difference between these two functions is in their parameter type. However, all pointer types can be considered equivalent since they can be bitcast in a lossless way. These functions are usually produced by copy-and-paste programming, where a given code pattern is copied and then repurposed [3, 28, 32]. The second pair of functions, shown in Figure 3.1b, were extracted from the `403.gcc` benchmark and they are fully identical. These functions are part of GCC’s backend, where it is common to have code that is automatically generated from a machine description [22, 34, 42].

<pre>void delete_contexts_MotionInfo (MotionInfoContexts *enco_ctx) { if(enco_ctx == NULL) return; free(enco_ctx); return; }</pre>	<pre>void delete_contexts_TextureInfo (TextureInfoContexts *enco_ctx) { if(enco_ctx == NULL) return; free(enco_ctx); return; }</pre>
--	--

(a) Two semantically identical functions extracted from the `482.sphinx3` benchmark.

<pre>rtx gen_floatitf2 (operand0, operand1) rtx operand0; rtx operand1; { return gen_rtx_SET (VOIDmode, operand0, gen_rtx_FLOAT (TFmode, operand1)); }</pre>	<pre>rtx gen_floatsitf2 (operand0, operand1) rtx operand0; rtx operand1; { return gen_rtx_SET (VOIDmode, operand0, gen_rtx_FLOAT (TFmode, operand1)); }</pre>
--	---

(b) Two semantically identical functions extracted from the `403.gcc` benchmark.

Figure 3.1: Example of identical functions.

<pre>template <int dim> unsigned PolynomialSpace<dim>:: compute_n_pols (const unsigned n) { unsigned n_pols = n; for (unsigned i=1; i<dim; ++i) { n_pols *= (n+i); n_pols /= (i+1); } return n_pols; }</pre>	<pre>template <int dim> inline unsigned TensorProductPolynomials<dim>:: x_to_the_dim (const unsigned x) { unsigned y = 1; for (unsigned d=0; d<dim; ++d) { y *= x; } return y; }</pre>
<p>----- After template specialization and applying optimizations: -----</p>	
<pre>unsigned PolynomialSpace<1>:: compute_n_pols(const unsigned n) { return n; }</pre>	<pre>unsigned TensorProductPolynomials<1>:: x_to_the_dim(const unsigned x) { return x; }</pre>

Figure 3.2: Two function extracted from the `447.dealIII` benchmark that are not identical at the source level, but after applying template specialisation and optimisations they become identical at the IR level.

Note, however, that functions can be identical at the IR or machine level without necessarily being identical at the source level. Figure 3.2 shows two real functions extract from the 447.dealIII program in the SPEC CPU2006 [55] benchmark suite. Although these two functions are not identical at the source level, they become identical after a template specialisation and some optimisations are applied, in particular, constant propagation, constant folding, and dead-code elimination. Specialising `dim` to 1 enables to completely remove the loop in the function `PolynomialSpace`. Similarly, specializing `dim` to 1 results in only the first iteration of the loop in the function `TensorProductPolynomials` being executed. The compiler is able to statically analyze and simplify the loops in both functions, resulting in the identical functions shown at the bottom of Figure 3.2.

Identical code is particularly common in C++ programs with heavy use of *parametric polymorphism*, via template or *auto* type deduction.

3.2.1 Merging Identical Object Code During Link Time

The simplest way of merging identical functions is by looking at their object code, during link time. *Identical code folding* (ICF) is an optimisation that identifies and merges two or more read-only sections, typically functions, that have identical contents. This optimisation is commonly found in major linkers, such as *gold* [35, 56], LLVM’s *lld*, and the MSVC linker [1].

Before applying ICS, it is a common practice for linkers to place functions in separate sections [35, 56]. Therefore, merging identical functions can be generalised to the problem of merging identical sections. In the simplest case, two functions or sections are identical if they have exactly the same binary representation. In general, two sections are considered identical if they have the identical section flags, data, code, and relocations. Two relocations are considered identical if they have the same relocation types, values, and if they point to the same or, recursively, identical sections.

Once a set of identical functions have been identified, merging them requires a simple operation. Only one of the functions in the set is kept for the final binary and all the other copies are discarded. Afterwards, every reference to the discarded functions must be redirected to the kept function.

Since the equality relation has a cyclic definition, ICF is defined as a fixed-point computation, i.e., it is applied repeatedly until a convergence is obtained. There are two approaches with distinct trade-offs: (i) The pessimistic approach starts with all

sections marked as being different and then repeatedly compares them trying to prove their equality, grouping those found to be identical, including their relocations. This approach is implemented in the widely used *gold* linker. (ii) The optimistic approach starts with all functions marked as potentially identical and then repeatedly compares them trying to disprove their equality, partitioning those found to be different. This approach is implemented in LLVM's linker, *lld*.

3.2.2 Identical Function Merging

A similar optimisation for merging identical functions, but instead at the intermediate representation (IR) level, is also offered by both GCC and LLVM [2, 39]. This optimisation is only flexible enough to accommodate simple type mismatches provided they can be bitcast in a lossless way.

A very strict function comparator is used to identify if two functions are semantically equivalent. First it compares the signature and other general attributes of the two functions. The functions must have identical signature, i.e., the same return type, the same number of arguments, and exactly the same list of argument types. Then this function comparator performs a simultaneous walk, in depth first order, in the functions' control-flow graphs. This walk starts at the entry block for both functions, then takes each block from each terminator in order. As an artifact, this also means that unreachable blocks are ignored. Finally, it iterates through each instruction in each basic block. Two blocks are equivalent if they have equivalent instructions in exactly the same order, without excess. The comparator always fails conservatively, erring on the side of claiming that two functions are different.

When a pair of equivalent functions is identified, we can create either an alias or a *thunk*. Aliasing entails eliminating one of the functions and replacing all its call-sites to the other function. Thunks must be created when neither of the equivalent functions can be eliminated by aliasing. In such case, a thunk is created for either one of the functions, replacing its body by a call to the other function, which allows all call-sites and name references to both functions to be preserved. Aliasing is preferred since it is cheaper and adds no runtime overhead. The appropriate merging is applied according to following rules:

- If the address of at least one function is not taken, alias can be used.
- If the function is part of COMDAT section that can be replaced, we must use thunk.

- If we create a thunk and none of functions is writeable, we can redirect calls instead.

Although very restrictive, this optimisation guarantees that any pair of mergeable functions will result in code size reduction with no performance overhead.

Its simplicity also allows for an efficient exploration approach based on computing a hash of the functions and then using a binary tree to identify equivalent functions. Since hashing is cheap to compute, it allows us to efficiently group possibly equivalent functions and filter out functions that are obviously unique. This hash must have the property that if function $F = G$ according to the comparison function, then $hash(F) = hash(G)$. Therefore, as an optimisation, two functions are only compared if they have the same hash. This consistency property is critical to ensuring all possible merging opportunities are exploited. Collisions in the hash affect the speed of the pass but not the correctness or determinism of the resulting transformation.

A function hash is calculated by considering only the number of arguments and whether a function contains variadic parameters, the order of basic blocks (given by the successors of each basic block in depth first order), and the order of opcodes of each instruction within each of these basic blocks. This mirrors the strategy of `compare()` uses to compare functions by walking the BBs in depth first order and comparing each instruction in sequence. Because this hash does not look at the operands, it is insensitive to things such as the target of calls and the constants used in the function, which makes it useful when possibly merging functions which are the same modulo constants and call targets.

All functions can be sorted based on their hash value, which ends up grouping possibly equivalent functions together. If the hash value of a given function matches any of its adjacent values in the sorted list, this function must be considered for merging. Functions with a unique hash value can be easily ignored since no other function will be found equivalent.

The functions that remain are inserted into a binary tree, where functions are the node values themselves. An order relation is defined over the set of functions. We need total-ordering, so we need to maintain four properties on the functions set:

- $a \leq a$ (reflexivity);
- if $a \leq b$ and $b \leq a$ then $a = b$ (antisymmetry);
- if $a \leq b$ and $b \leq c$ then $a \leq c$ (transitivity);

- for all a and b , $a \leq b$ or $b \leq a$ (totality).

This total-ordering was made through special function comparison procedure that returns:

- 0 when functions are semantically equal,
- -1 when the left hand function is less than the right hand function, and
- 1 for the opposite case.

Functions are kept on binary tree. For each new function F we perform lookup in binary tree.

3.3 Merging Beyond Identical Functions

In the previous sections, we have seen compiler optimisations that merge identical functions. However, nearly identical functions, with only minor differences, are also commonly found. Figure 3.3 shows two examples of nearly identical functions found in real programs. The highlighted differences prevent these functions from being merged by the identical function merging techniques. The first pair of functions, shown in Figure 3.3a, illustrates code that is usually produced by copy-and-paste programming, where a given code pattern is copied and then repurposed [3, 28, 32]. The second pair of functions, shown in Figure 3.3b, are produced by generative programming [13, 16], where their code was automatically generated using a description language [22].

Edler von Koch et al. [19] have proposed a function-merging technique which exploits structural similarity among functions. Their optimization is able to merge nearly identical functions. Two functions are structurally similar if both their function types are equivalent and their CFGs are isomorphic. Two function types are equivalent if they agree in the number, order, and types of their parameters as well as their return types, linkage type, and other compiler-specific properties. In addition to the structural similarity of the functions, their technique also requires that corresponding basic blocks have exactly the same number of instructions and that corresponding instructions must have equivalent resulting types. Mergeable functions are only allowed to differ in corresponding instructions, where they can differ in their opcodes or in the number and type of their input operands. Corresponding named values must have the same data type.

```

double DMin(double *vec, int n) {
    int i;
    double best;
    best = vec[0];
    for (i=1; i<n; i++)
        if (vec[i] < best)
            best = vec[i];
    return best;
}

double DMax(double *vec, int n) {
    int i;
    double best;
    best = vec[0];
    for (i=1; i<n; i++)
        if (vec[i] > best)
            best = vec[i];
    return best;
}

```

(a) Two similar functions extracted from the 456.hammer benchmark.

```

rtx
gen_peephole2_1270 (curr_insn, operands)
    rtx curr_insn ATTRIBUTE_UNUSED;
    rtx *operands;
{
    rtx_val = 0;
    HARD_REG_SET _regs_allocated;
    CLEAR_HARD_REG_SET (_regs_allocated);
    start_sequence ();
    operands[2] = GEN_INT(
        exact_log2(INTVAL(operands[1])));
    emit (gen_rtx_PARALLEL (VOIDmode,
        gen_rtxvec (2,
            gen_rtx_SET (VOIDmode,
                operands[0],
                gen_rtx_ASHIFT (SImode,
                    copy_rtx (operands[0]),
                    operands[2])),
            gen_rtx_CLOBBER (VOIDmode,
                gen_rtx_REG (CCmode,
                    17))));
    _val = gen_sequence ();
    end_sequence ();
    return _val;
}

rtx
gen_peephole2_1271 (curr_insn, operands)
    rtx curr_insn ATTRIBUTE_UNUSED;
    rtx *operands;
{
    rtx_val = 0;
    HARD_REG_SET _regs_allocated;
    CLEAR_HARD_REG_SET (_regs_allocated);
    start_sequence ();
    operands[2] = GEN_INT(
        exact_log2(INTVAL(operands[1])));
    emit (gen_rtx_PARALLEL (VOIDmode,
        gen_rtxvec (2,
            gen_rtx_SET (VOIDmode,
                operands[0],
                gen_rtx_ASHIFT (DImode,
                    copy_rtx (operands[0]),
                    operands[2])),
            gen_rtx_CLOBBER (VOIDmode,
                gen_rtx_REG (CCmode,
                    17))));
    _val = gen_sequence ();
    end_sequence ();
    return _val;
}

```

(b) Two similar functions extracted from the 403.gcc benchmark.

Figure 3.3: Example of two pairs of highly similar functions. Because they are not identical, they cannot be merged by the function merging technique currently found in major compilers.

Because their technique is limited to functions with identical CFGs and function types, once it merges a pair of functions, a third *similar* function cannot be merged into the resulting merged function since they will differ in both CFGs and their lists of parameters. Due to this limiting factor, the state-of-the-art has to first group mergeable functions before simultaneously merging all functions within a group.

Their algorithm iterates simultaneously over corresponding basic blocks in the set functions being merged, as they have isomorphic CFGs. Figure 3.4 shows an example of two functions with isomorphic CFGs and their corresponding basic blocks arranged side by side.

Every pair of basic blocks have their instructions analysed in a pairwise manner.

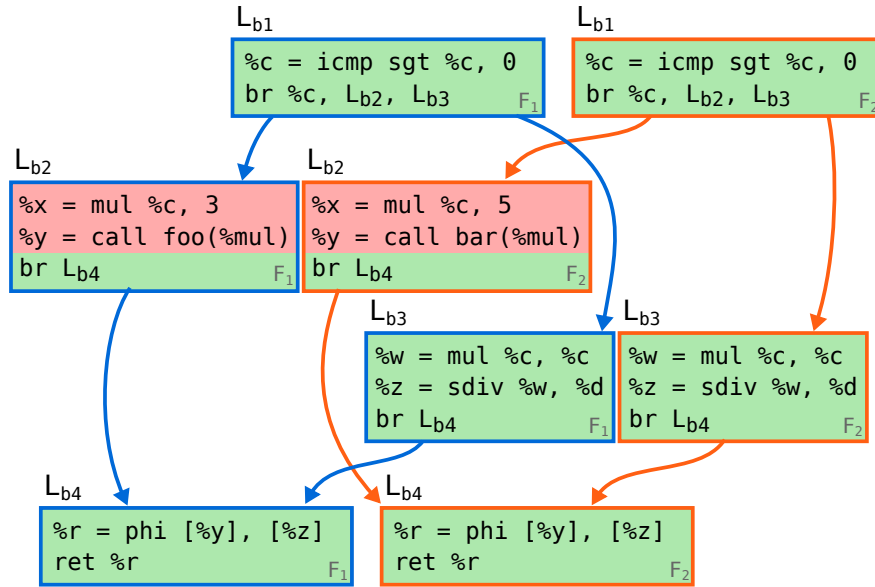


Figure 3.4: An example of two functions with isomorphic CFGs and their corresponding basic blocks arranged side by side. Instructions in paired basic blocks are compared in a pairwise manner.

Two instructions match if they have the same opcode with equivalent data types and operands. Even if two instructions differ only on their operands, they are classified as mismatching. For every pair of basic blocks, if their corresponding instructions have any difference, except for the data type of the computed value, the merged basic block is split by inserting a switch branch to select which instruction to execute depending on a function identifier. A phi-node is used to unify the mismatching instructions as a single named value. This unification is only possible because they compute values of the equivalent data types. Note that no operand selection is performed, every use of the mismatching instructions will refer to their phi-node. Figure 3.5 shows an example of a merged basic block containing two mismatching pairs of instructions. A split is added for every pair of mismatching instructions with the phi-node instruction added to the their immediate point of convergence.

Overall, except for mismatching pairs of instructions, the two functions must have identical function types, i.e., they must have the same return type and list of arguments, identical CFGs, with corresponding basic blocks having the same number of instructions. Although this technique improves over LLVM’s identical function merging, it is still unnecessarily limited. In Section 6.1, we showed examples of very similar real functions where the technique proposed by Edler von Koch et al. [19] fails to merge. In Chapter 4 we introduce a novel technique that addresses such limitations improving

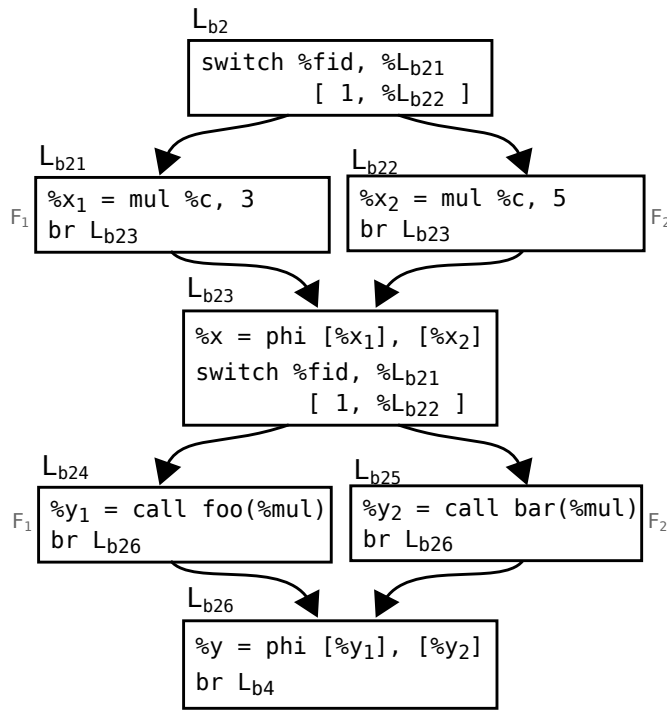


Figure 3.5: An example of a merged basic block containing two mismatching pairs of instructions. A split is added for every pair of mismatching instructions with the phi-node instruction added to the their immediate point of convergence.

on their technique across the board.

3.4 Code Factoring

Code factoring refers to related techniques that address the same fundamental problem of duplicated code in a different way. Code factoring can be applied at different levels of the program [40]. Local factoring, also known as local code motion, moves identical instructions from multiple basic blocks to either their common predecessor or successor, whenever valid [5, 33, 40]. Procedural abstraction finds identical code that can be extracted into a separate function, replacing all replicated occurrences with a function call [17, 40].

Procedural abstraction differs from function merging as it usually works on single basic blocks or single-entry single-exit regions. Moreover, it only works for identical segments of code, and every identical segment of code is extracted into a separate new function. Function merging, on the other hand, works on whole functions, which can be identical or just partially similar, producing a single merged function.

However, all these techniques are orthogonal to the proposed optimization and

could complement each other at different stages of the compilation pipeline.

3.5 Code Similarity

Code similarity has also been used in other compiler optimizations or tools for software development and maintenance. In this section, we describe some of these applications.

Coutinho et al. [10] proposed an optimization that uses instruction alignment to reduce divergent code for GPU. They are able to fuse divergent branches that contain single basic blocks, improving GPU utilization.

Similarly, analogous algorithms have also been suggested to identify the differences between two programs, helping developers with source-code management and maintenance [41, 64]. These techniques are applied in tools for source-code management, such as the *diff* command [41].

Similar techniques have also been applied to code editors and IDEs [51, 58]. For example, SourcererCC [51] detects possible clones, at the source level, by dividing the programs into a set of code blocks where each code block is itself represented by a bag-of-tokens, i.e., a set of tokens and their frequencies. Tokens are keywords, literals, and identifiers, but not operators. Code blocks are considered clones if their degree of similarity is higher than a given threshold. In order to reduce the number of blocks compared, candidate blocks are filtered based on a few of their tokens where at least one must match.

Our ranking mechanism uses an approach similar to SourcererCC, where we use opcode frequencies and type frequencies to determine if two functions are likely to have similar code. However, we need a precise and effective analysis of code similarity when performing the actual merge. To this end, we use a sequence alignment technique.

Chapter 4

Function Merging by Sequence Alignment

In this chapter, we present our novel function merging technique, called FMSA. Our approach is based upon the concept of sequence alignment, developed in bioinformatics for identifying functional or evolutionary relationships between different DNA or RNA sequences. Similarly, we use sequence alignment to find areas of functional similarity in arbitrary function pairs. Aligned segments with equivalent code are merged. The remaining segments where the two functions differ are added to the new function too but have their code guarded by a function identifier. This approach leads to significant code size reduction.

Applying sequence alignment to all pairs of functions is prohibitively expensive even for medium sized programs. To counter this, our technique is integrated with a ranking-based exploration mechanism that efficiently focuses the search to the most promising pairs of functions. As a result, we achieve our code size savings while introducing little compilation-time overhead.

Compared to identical function merging, we introduce extra code to be executed, namely the code that chooses between dissimilar sequences in merged functions. A naive implementation could easily hurt performance, e.g by merging two hot functions with only few similarities. Our implementation can avoid this by incorporating profiling information to identify blocks of hot code and effectively minimize the overhead in this portion of the code.

In this chapter, we make the following contributions:

- We are the first to allow merging arbitrary functions, even ones with different signatures and CFGs.

- A novel ranking mechanism for focusing inter-procedural optimizations to the most profitable function pairs.
- Our function merging by sequence alignment technique is able to reduce code size by up to 25% on Intel and 30% on ARM, significantly outperforming the previous state of the art, proposed by Edler von Koch et al. [19], while introducing minimal compile-time and negligible run-time overheads.

4.1 Motivation

In this section we make the argument for a more powerful function merging approach. Consider the examples from two SPEC CPU2006 benchmarks shown in Figures 4.1 and 4.2.

```

glist_t glist_add_float32(glist_t g, float32 val){
    gnode_t *gn;
    gn = (gnode_t *) mymalloc (sizeof(gnode_t));
    gn->data.float32 = val;
    gn->next = g;
    return ((glist_t) gn);
}

glist_t glist_add_float64(glist_t g, float64 val){
    gnode_t *gn;
    gn = (gnode_t *) mymalloc (sizeof(gnode_t));
    gn->data.float64 = val;
    gn->next = g;
    return ((glist_t) gn);
}
-----Merged Function-----
glist_t merged(bool func_id,
               glist_t g, float32 v32, float64 v64){
    gnode_t *gn;
    gn = (gnode_t *) mymalloc (sizeof(gnode_t));
    if (func_id)
        gn->data.float32 = v32;
    else
        gn->data.float64 = v64;
    gn->next = g;
    return ((glist_t) gn);
}

```

Figure 4.1: Example of two functions from the benchmark `sphinx` with different parameters that could be merged, as shown at the bottom. We highlight where they differ.

Figure 4.1 shows two functions from the `482.sphinx3` benchmark. The two functions are almost identical, only their function arguments are of different types, `float32` and `float64`, causing a single operation to be different. As shown at the bottom of

Figure 4.1, these functions can be easily merged in three steps. First, we expand the function argument list to include the two parameters of different types. Then, we add a function identifier, `func_id`, to indicate which of the two functions is called. Finally, we place the lines that are unique to one of the functions in a conditional branch predicated by the `func_id`. Overall, merging these two functions reduces the total number of machine instructions by 18% in the final object file for the Intel x86 architecture.

Despite being so similar, neither GCC nor LLVM can merge the two functions. They can only handle identical functions, allowing only for type mismatches that can be removed by lossless bitcasting of the conflicting values. Similarly, the previous state of the art [19], while more powerful, cannot merge the two functions either. It requires both functions to have the same list of parameters.

```

void q_cond_phase(
int control, int target, qreg *reg){
    int i;
    COMPLEX_FLOAT z;
    if(q_objc_put(COND_PHASE, control, target))
        return;
    z = q_cexp(pi / (1 << (control - target)));
    for(i=0; i<reg->size; i++) {
        if(reg->node[i].state & (1 << control)) {
            if(reg->node[i].state & (1 << target))
                reg->node[i].amplitude *= z;
        }
    }
    q_decohere(reg);
}

void q_cond_phase_inv(
int control, int target, qreg *reg){
    int i;
    COMPLEX_FLOAT z;
    z = q_cexp(-pi / (1 << (control - target)));
    for(i=0; i<reg->size; i++) {
        if(reg->node[i].state & (1 << control)) {
            if(reg->node[i].state & (1 << target))
                reg->node[i].amplitude *= z;
        }
    }
    q_decohere(reg);
}

```

(a) Pair of similar input functions.

```

void merged(bool func_id,
int control, int target, quantum_reg *reg){
    int i;
    COMPLEX_FLOAT z;
    if(func_id)
        if(quantum_objcode_put(COND_PHASE, control, target))
            return;
    float var = (func_id)?pi:(-pi);
    z = quantum_cexp(var / (1 << (control - target)));
    for(i=0; i<reg->size; i++) {
        if(reg->node[i].state & (1 << control)) {
            if(reg->node[i].state & (1 << target))
                reg->node[i].amplitude *= z;
        }
    }
    quantum_decohere(reg);
}

```

(b) Function produced by merging both input functions.

Figure 4.2: Example of two functions from the benchmark `libquantum` with different CFGs that could be merged, as shown at the bottom. We highlight where they differ.

Figure 4.2 shows another two functions extracted from `462.libquantum`. While these two functions have the same signature, i.e. the same return type and list of parameters, they differ slightly in their bodies. Merging them manually is straightforward, as shown at the bottom of Figure 4.2, reducing the number of instructions by 23% in the final object file. But again, none of the existing techniques can merge the two functions. The state-of-the-art can work with non-identical functions, but it needs their CFGs to be identical. Even a single extra basic block, as in this case, makes merging impossible.

These examples show that all existing techniques are severely limited. Optimization passes in production compilers work only on effectively identical functions. State-of-the-art techniques can merge functions only when they are structurally identical, with isomorphic CFGs, and identical signatures. All of them miss massive opportunities for code size reduction. In the next sections, we show a better approach which removes such constraints and is able to merge arbitrary functions, when it is profitable to do so.

4.2 Our Approach

In this section we describe our proposed function merging technique and show how it merges the motivating examples. Our technique works on any two arbitrary functions, even when they have few similarities and merging them would be counter-productive. For that reason, we also introduce a cost model to decide when it is beneficial to merge two functions (see Section 4.3.1). To avoid an expensive quadratic exploration, we integrate our profitability analysis with an efficient ranking mechanism based on a lightweight fingerprint of the functions.

4.2.1 Overview

Intuitively, when we are manually merging two functions, in a textual format, we try to visualize them side by side, identifying the equivalent segments of code and the non-equivalent ones. Then, we use this understanding to create the merged function. In this chapter, we propose a technique that follows this simple yet effective principle. At the core of our technique lies a sequence alignment algorithm, which is responsible for arranging the code in segments that are either equivalent or non-equivalent. We implement this technique at the level of the intermediate representation (IR). Our current

implementation assumes that the input functions have all their ϕ -functions demoted to memory operations, simplifying our code generation.




Figure 4.3: Overview of our function-merging technique. Equivalent segments of code is represented in light green and the non-equivalent ones in dark red.

The proposed technique consists of three major steps, as depicted in Figure 4.3. First, we linearize each function, representing the CFG as a sequence of labels and instructions. The second step consists of applying a sequence alignment algorithm, borrowed from bioinformatics, which identifies regions of similarity between sequences. The sequence alignment algorithm allows us to arrange two linearized functions into segments that are equivalent between the two functions and segments where they differ from one another. The final step performs the code generation, actually merging the two functions. Aligned segments with equivalent code are merged, avoiding redundancy, while the remaining segments where the two functions differ have their code guarded by a function identifier.

During code generation, we create a merged list of parameters, including the extra function identifier if there are any dissimilar segments. Arguments of the same type are shared between the functions, without necessarily keeping their original order. We also allow for the return types to be different, in which case we use an aggregate type to return values of both types. If one of them is void, then we do not create an aggregate type, we just return the non-void type. Given the appropriate function identifier, the merged function is semantically equivalent to the original functions, so we replace all of their invocations with the new function. It should be noted that in the special case where we merge identical functions, the output is also identical, emulating the behavior of function merging in production compilers.

After producing the merged function, the bodies of the original functions are replaced by a single call to this new function, creating what is sometimes called a *thunk*. In some cases, it may also be valid and profitable to completely delete the original functions, remapping all their original calls to the merged function. Two of the key

facts that prohibit the complete removal of the original functions are the existence of indirect calls or the possibility of external linkage.

4.2.2 Linearization




Figure 4.4: Linearizing the CFG of an example function.

Linearization is a key step for enabling the use of sequence alignment. Although linearization of CFGs usually refers to a predicated representation, resulting from an if-conversion, in this thesis, we use a simpler definition. Our linearization takes the CFG of the function, specifies a traversal order of the basic blocks, and for each block outputs its label and its instructions. It maintains the original ordering of the instructions inside each basic block. The edges of the CFG are implicitly represented with branch instructions having the target labels as operands. Figure 4.4 shows a simplified example of linearizing the CFG of a real function extracted from the SPEC CPU2006 400.perlbench benchmark.

The traversal order we use for linearization has no effect on the correctness of the transformation but it can impact its effectiveness. We empirically chose a reverse post-order traversal with a canonical ordering of successor basic blocks. This strategy leads to good performance in our experiments.

4.2.3 Sequence Alignment

When merging two functions, the goal is to identify which pairs of instructions and labels that can be merged and which ones need to be selected based on the actual function being executed. To avoid breaking the semantics of the original program, we

also need to maintain the same order of execution of the instructions for each one of the functions.

To this end, after linearisation, we reduce the problem of merging functions to the problem of *sequence alignment*. Figure 4.5 shows an example of the sequence alignment between two linearised functions extracted from the `400.perlbench` benchmark in SPEC CPU2006 [55].




Figure 4.5: The sequence alignment between two functions.

Specifically for the function merging, we are concerned with the alphabet consisting of all possible typed instructions and labels. Every linearised function represents a sequence derived from this alphabet. We explain the equivalence relation used for this alphabet in the next section. Although we only consider pair-wise alignments, the technique would also work for multi-sequences.

Our work uses the Needleman-Wunsch algorithm [44], with the standard scoring scheme, to perform sequence alignment. This algorithm gives an alignment that is guaranteed to be optimal for a given scoring scheme [27], however, other algorithms could also be used with different performance and memory usage trade-offs [7, 26, 44, 54]. Different alignments would produce different but valid merged functions.

4.2.4 Equivalence Evaluation

Before we merge functions, we first need to define what makes two pieces of code equivalent and therefore mergeable. In this section, we define equivalence in two separate cases, the equivalence between instructions and the equivalence between labels.

In general, two instructions are equivalent if: (1) their opcodes are semantically equivalent, but not necessarily the same; (2) they both have equivalent types; and (3) they have pairwise operands with equivalent types. Types are equivalent if they can be

bitcast in a lossless way from one to the other. For pointers, we also need to make sure that there is no conflict regarding memory alignment. In the special case of function calls, type equivalence means that both instructions have identical function types, i.e. identical return types and identical list of parameters.

Labels can represent both normal basic blocks and landing blocks used in exception handling code. Labels of normal basic blocks are ignored during code equivalence evaluation, but we cannot do the same for landing blocks. We describe how we handle such blocks in more detail in the following section.

Exception Handling Code

Most modern compilers, including GCC and LLVM, implement the zero-cost Itanium ABI for exception handling [14] sometimes called the *landing-pad* model. This model consists of: (1) invoke instructions that have two successors, one for the normal execution and one for handling exceptions, called the landing block; (2) landing-pad instructions that encode which action is taken when an exception has been thrown. The invoke instruction co-operates tightly with its landing block. The landing block must have a landing-pad instruction as its first non- ϕ instruction. As a result, two equivalent invoke instructions must also have landing blocks with identical landing-pad instructions. This verification is made easy by having the landing-pad instruction as the first instruction in a landing block. Similarly, landing-pad instructions are equivalent if they have exactly the same type and also encode identical lists of exception and cleanup handlers.

4.2.5 Code Generation

The code generation phase is responsible for producing a new function from the output of the sequence alignment. Our four main objectives are: merging the parameter lists; merging the return types; generating select instructions to choose the appropriate operands in merged instructions; and constructing the CFG of the merged function.

Our approach can effectively handle multiple different function merging scenarios:

- identical functions,
- functions with differing bodies,
- functions with different parameter lists,

- functions with different return types,
- and any combination of these cases.

To maintain the semantics of the original functions, we must be able to pass their parameters to the new merged function. The merged parameter list is the union of the original lists, with placeholders of the correct type for any of the parameters. Maintaining the original order is not important for maintaining semantics, so we make no effort to do so. If the two functions have differing bodies, we add an extra binary parameter, called the function identifier, to the merged list of parameters. This extra parameter is required for selecting code that should be executed only for one of the merged functions.

Figure 4.6 depicts how we merge the list of parameters of two functions. First, we create the binary parameter that represents the function identifier, one of the functions will be identified by the value `true` and the other by the value `false`. We then add all the parameters of one of the functions to the new list of parameters. Finally, for each parameter of the second function, we either reuse an existing and available parameter of identical type from the first function or we add a new parameter. We keep track of the mapping between the lists of parameters of the original functions and the merged function so that, later, we are able to update the function calls. When replacing the function calls to the new merged function, parameters that are not used by the original function being called will receive undefined values.

The reuse of parameters between the two merged functions provides the following benefits: (1) it reduces the overheads associated with function call abstractions, such as reducing the number of values required to be communicated between functions. (2) if both functions use merged parameters in similar ways, it will remove some of the cases where we need select instructions to distinguish between the functions.




Figure 4.6: Example of a merge operation on the parameter lists of two functions.

There are multiple valid ways of merging parameter lists. For example, multiple parameters of one function may have the same type as a given parameter from the other

function. In such cases, we select parameter pairs that minimize the number of select instructions. We find them by analyzing all pairs of equivalent instruction that use the parameters as operands. Our experiments show that maximizing the matching of parameters, compared to never merging them, improves code-size reduction of individual benchmarks by up to 7%.

Our technique is able to merge any return types. When merging return types, we select the largest one as the base type. Then, we use bitcast instructions to convert between the types. Before a return instruction, we bitcast the values to the base return type. We reverse this at the call-site, where we cast back to the original type. Having identical types or void return are just special cases where casting is unnecessary. In the case of void types, we can just return *undefined* values since they will be discarded at the corresponding call-sites.

After generating the merged list of parameters and return type, we produce the CFG of the merged function in two passes over the aligned sequence. The first pass creates the basic blocks and instructions. The second assigns the correct operands to the instructions and connects the basic blocks. A two-passes approach is required in order to handle loops, due to cyclic data dependencies.

First, for each entry in the aligned sequence, we either create a new basic block for labels or we add a cloned instruction to the appropriate basic block. If the label represents a landing block, a landing-pad instruction is also added to the new basic block. During this process, we keep a mapping from the instructions and labels in the original functions to their corresponding values in the new merged function. We need this mapping to generate the use-definition chains for the merged function, which is done by pointing the operands of the instructions to the correct values in the function. However, at this point, the cloned instructions are given empty operands, as we are still creating the complete mapping.

While iterating over the aligned sequence, we also need to create extra basic blocks and branch instructions in order to maintain the semantics of the original functions, guarding the execution of instructions that are unique to one of the functions being merged. When transitioning from matching instructions or labels to non-matching ones, we need to branch to new basic blocks based on the function identifier. When transitioning back from non-matching segments to a matching segment, we need to reconnect both divergent points by branching back to a single new basic block where merged instructions will be added. This process generates diamond shaped structures in the CFG.

The second pass over the aligned sequence creates the operands of all instructions. We use the previously created mapping in order to identify the correct operands for each instruction in the merged function. There are two main cases: (1) Creating the operands for non-matching instructions (i.e. those that occur in just one function) is straightforward. In this case, we only need to use the values on which the operands of the original instruction map. (2) Matching instructions can have different values in corresponding operands in each one of the original functions. If this is the case and the original operands map to different values V_1 and V_2 , then we need to choose at runtime the correct value based on the function identifier. We do with an extra select instruction “`select (func_id==1), V_1, V_2`”, which computes the operand of the merged instruction. If the two values are statically identical, then we do not need a select.

If the operands are labels, instead of adding a select instruction, we perform operand selection through divergent control flow, using a new basic block and a conditional branch on the function identifier. If the two labels represent landing blocks, we hoist the landing-pad instruction to the new common basic block, converting it to a landing block and converting the two landing blocks to normal basic blocks. This is required for the correctness of the landing-pad model.

Similar to previous work on vectorization [48], we also exploit commutative instructions in order to maximize similarity. When assigning operands to commutative instructions, we perform operand reordering to maximize the number of matching operands and reduce the total number of select instructions required. It is also important to note that if we are merging two identical functions, no select or extra branch instruction will be added. As a result, we can remove the extra parameter that represents the function identifier.

4.3 Focusing on Profitable Functions

Although the proposed technique is able to merge any two functions, it is not always profitable to merge them. In fact, as it is only profitable to merge functions that are sufficiently similar, for most pairs of functions, merging them increases code size. In this section, we introduce our framework for efficiently exploring the optimization space, focusing on pairs of functions that are profitable to merge.

For every function, ideally, we would like to try to merge it with all other functions and choose the pair that maximizes the reduction in code size. However, this quadratic

exploration over all pairs of functions results in prohibitively expensive compilation overhead. In order to avoid the quadratic exploration of all possible merges, we propose the exploration framework shown in Figure 4.7 for our optimization.




Figure 4.7: Overview of our exploration framework.

The proposed framework is based on a light-weight ranking infrastructure that uses a *fingerprint* of the functions to evaluate their similarity. It starts by precomputing and caching fingerprints for all functions. The purpose of fingerprints is to make it easy to discard unpromising pairs of functions so that we perform the more expensive evaluation only on the most promising pairs. To this end, the fingerprint consists of: (1) a map of instruction opcodes to their frequency in the function; (2) the set of types manipulated by the function. While functions can have several thousands of instructions, an IR usually has just a few tens of opcodes, e.g., the LLVM IR has only about 64 different opcodes. This means that the fingerprint needs to store just a small integer array of the opcode frequencies and a set of types, which allows for an efficient similarity comparison.

By comparing the opcode frequencies of two functions, we are able to estimate the best case merge, which would happen if all instructions with the same opcode could match. This is a very optimistic estimation. It would be possible only if instruction types and order did not matter. We refine it further by estimating another best case merge, this time based on type frequencies, which would happen if all instructions

with the same data type could match.

Therefore, the upper-bound reduction, computed as a ratio, can be generally defined as

$$UB(f_1, f_2, K) = \frac{\sum_{k \in K} \min\{freq(k, f_1), freq(k, f_2)\}}{\sum_{k \in K} freq(k, f_1) + freq(k, f_2)}$$

where $UB(f_1, f_2, Opcodes)$ computes the opcode-based upper bound and $UB(f_1, f_2, Types)$ computes the type-based upper bound. The final estimate selects the minimum upper bound between the two, i.e.,

$$s(f_1, f_2) = \min\{UB(f_1, f_2, Opcodes), UB(f_1, f_2, Types)\}$$

This estimate results in a value in the range $[0, 0.5]$, which encodes a description that favours maximizing both the opcode and type similarities, while also minimizing their respective differences. Identical functions will always result in the maximum value of 0.5.

For each function f_1 , we use a priority queue to rank the topmost similar candidates based on their similarity, defined by $s(f_1, f_2)$, for all other functions f_2 . We use an exploration threshold to limit how many top candidates we will evaluate for any given function. We then perform this candidate exploration in a greedy fashion, terminating after finding the first candidate that results in a profitable merge and committing that merge operation.




Figure 4.8: Average CDF for the position of the profitable candidate and the percentage of merged operations covered. 89% of the merge operations happen with the topmost candidate.

Ideally, profitable candidates will be as close to the top of the rank as possible. Figure 4.8 shows the cumulative distribution of the position of the profitable candidates in a top 10 rank. It shows that about 89% of the merge operations occurred with the

topmost candidate, while the top 5 cover over 98% of the profitable candidates. These results suggest that fingerprint similarity is able to accurately capture the real function similarity, while reducing the exploration cost by orders of magnitudes, depending on the actual number and size of the functions.

When a profitable candidate is found, we first replace the body of the two original functions to a single call to the merged function. Afterwards, if the original functions can be completely removed, we update the call graph, replacing the calls to the original functions by calls to the merged function. Finally, the new function is added to the optimization working list. Because of this feedback loop, merge operations can also be performed on functions that resulted from previous merge operations.

4.3.1 Profitability Cost Model

Although the proposed technique is able to merge any two functions, it is not always profitable to do so. In fact, as it is only profitable to merge functions that are sufficiently similar, for most pairs of functions, merging them increases code size. Therefore, we must be able to decide which merge operations are profitable.

A merge operation is profitable when replacing the original pair of functions by the merged function results in an overall smaller code. Since merging most pairs of functions tend to produce a much larger function, the profitability analysis is critical for the optimisation strategy.

In order to estimate the code-size benefit, we first estimate the size of all three functions. The size of each function is estimated by summing up the estimated binary size of all instruction in the function. The binary size of each instruction is estimated by querying the compiler's built-in target-specific cost model. As described in Section 2.1.2, although the cost model estimations are not necessarily precise, they are useful for comparing two pieces of code and deciding when a replacement is profitable.

In addition to measuring the difference in size of the merged function, we also need to take into account all extra costs involved in replacing the original functions with the merged one: (1) for the cases where we need to keep a *thunk* of the original functions, preserving their original signatures; and (2) for the cases where we update the call graph, there might be an extra cost with a call to the merged function due to the increased number of arguments.

Let $c(f)$ be the code-size cost of a given function f , and $\delta(f_i, f_j)$ represent the extra

costs involved when replacing or updating function f_i with the function f_j . Therefore, given a pair of functions $\{f_1, f_2\}$ and the merged function $f_{1,2}$, we want to maximize the profit defined as:

$$\Delta(\{f_1, f_2\}, f_{1,2}) = (c(f_1) + c(f_2)) - (c(f_{1,2}) + \epsilon)$$

where $\epsilon = \delta(f_1, f_{1,2}) + \delta(f_2, f_{1,2})$. We consider that the merge operation is profitable if $\Delta(\{f_1, f_2\}, f_{1,2}) > 0$.

Even though the use of cost models offer a good trade-off between compilation time and accuracy, they are expected to contain inaccuracies. Because we are operating on the IR level, one IR instruction does not necessarily translate to one machine instruction. Several number of optimisations and code transformations will run afterwards, modifying the code. Moreover, we cannot know exactly how each IR instruction will be lowered without actually running the compiler's backend. The same IR instruction can be lowered to different machine instructions, depending on its surrounding context, the instruction selection algorithm, and many other factors. Therefore, there is also an inherent limitation of estimating the cost of each instruction separately of its context.

4.3.2 Link-Time Optimization

There are different ways of applying this optimization, with different trade-offs. We can apply our optimization on a per compilation-unit basis, which usually results in lower compilation-time overheads because only a small part of the whole program is being considered at each moment. However, this also limits the optimization opportunities, since only pairs of functions within the same translation unit would be merged.

On the other hand, our optimization can also be applied in the whole program, for example, during link-time optimization (LTO). Optimizing the whole program is beneficial for the simple fact that the optimization will have more functions at its disposal. It allows us to merge functions across modules.




Figure 4.9: In our experiments we use a compilation pipeline with a monolithic link-time optimization (LTO).

In addition to the benefit of being able to merge more functions, when optimizing the whole program, we can also be more aggressive when removing the original functions, since we know that there will be no external reference to them. However, if the optimization is applied per translation unit, then extra conditions must be guaranteed, e.g., the function must be explicitly defined as internal or private to the translation unit.

Figure 4.9 shows an overview of the compilation pipeline used throughout our evaluation. First, we apply early code-size optimizations (`-Os`) to each compilation unit. Then, function merging and further code-size optimizations are applied during monolithic link-time optimization (LTO). With LTO, object file generation is delayed until all input modules are known, instead of being generated per translation unit, which enables more powerful optimizations based on whole-program analyses.

4.4 Evaluation

In this section, we evaluate the proposed optimization, where we analyze our improvements on code size reduction, as well as its impact on the program’s performance and compilation-time.

4.4.1 Experimental Setup

We compare our optimization against the state-of-the-art [19] and LLVM’s identical [2] function merging techniques. In our evaluation, we refer to the identical function merging as *Identical*, the state-of-the-art as *SOA*, and our approach as *FMSA*. We also run LLVM’s identical function merging before both *SOA* and *FMSA*, as this helps to reduce compilation time by efficiently reducing the number of trivially mergeable functions.

All optimizations are implemented in LLVM v8 and evaluated on two benchmark suites: the C/C++ SPEC CPU2006 [55] and MiBench [23]. We target two different instruction sets, the Intel x86-64 and the ARM Thumb. Our Intel test bench has a quad-core 3.4 GHz Intel Core i7 CPU with 16 GiB of RAM. The ARM test bench has a Cortex-A53 ARMv8 CPU of 1.4 GHz with 1 GiB of RAM. We use the Intel platform for compiling for either target. As a result, compilation-time is almost identical for both targets. Changing the target only affects the behavior of the backend, a very short part of the pipeline. Because of that, we only report compilation-time overhead results for one of the targets, the Intel ISA.

For the proposed optimization, we vary the exploration threshold (Section 4.3) and

we present the results for a range of threshold values. We also show the results for the oracle exploration strategy, which instead of using a rank-based greedy approach, merges a function with all candidates and chooses the best one. This oracle is a perfect ranking strategy but is unrealistic. It requires a very costly quadratic exploration, as explained in Section 4.3.

4.4.2 Code-Size Reduction

Figure 4.10 reports the code size reduction over the baseline for the linked object. We observe similar trends of code size reduction on both target architectures. This is expected because the optimizations are applied at the platform-independent IR level. Changing the target architecture introduces only second order effects, such as slightly different decisions due to the different cost model (LLVM’s TTI) and differences in how the IR is encoded into binary.




Figure 4.10: Object file size reduction for Intel (top) and ARM (bottom). We evaluate our approach (FMSA) under four different exploration thresholds, which control how many potential merging pairs we examine for each function before making a decision. Even for a threshold of one, we outperform the state-of-the-art by $2.4\times$ (Intel) and $1.9\times$ (ARM).

Our approach, FMSA, significantly improves over the state-of-the-art (SOA). For the Intel platform, FMSA can achieve an average code size reduction of up to 6.3% (or 6% with the lowest exploration threshold), while the SOA and Identical had an average reduction of 2.5% and 1.4%, respectively. Similarly, for the ARM platform, FMSA can achieve an average code size reduction of up to 6.1% (or 5.7% with the lowest threshold), while SOA and Identical had an average reduction of 3% and 1.8%, respectively.

For several of the benchmarks, the proposed technique achieves impressive code size reduction compared to other merging approaches.

In most cases, LLVM’s identical function merging has very little impact on code size. We see noticeable impact only on some of the C++ benchmarks, namely, `447.dealII`, `450.soplex`, `471.omnetpp`, `483.xalancbmk`. These are the cases that identical function merging was designed to handle, duplicate functions due to heavy use of templating. Although the state-of-the-art improves over LLVM’s identical function merging, it still gets most of its code size reduction for benchmarks with heavy use of templating. In addition to achieving better results in all of these cases, our technique also shows remarkable reductions on several of the C benchmarks, especially `462.libquantum` and `482.sphinx3`, where other techniques have no real impact.

In Section 6.1, we show two examples extracted from `462.libquantum` and `482.sphinx3`, where we detail how existing techniques fail to merge similar functions in these benchmarks. Our technique is the *first* that can handle these examples, producing merged functions equivalent to the handwritten ones shown in Figures 4.1 and 4.2.

Table 4.1: Number and size of functions present in each SPEC CPU2006 benchmark just before function merging, as well as number of merge operations applied by each technique.

Benchmarks	#Fns	Min/Avg/Max Size	Identical	SOA	FMSA[t=1]	FMSA[t=10]
400.perlbench	1699	1 / 125 / 12501	12	103	175	200
401.bzip2	74	1 / 206 / 5997	0	0	7	7
403.gcc	4541	1 / 127.7 / 20688	136	341	614	710
429.mcf	24	18 / 87.25 / 297	0	1	1	1
433.milc	235	1 / 67.69 / 416	0	6	26	34
444.namd	99	1 / 570.64 / 1698	1	1	5	5
445.gobmk	2511	1 / 43.22 / 3140	183	485	436	605
447.dealII	7380	1 / 60.63 / 4856	1835	2785	2974	3315
450.soplex	1035	1 / 73.27 / 1719	27	125	156	163
453.povray	1585	1 / 98.05 / 5324	60	112	193	212
456.hmmer	487	1 / 99.98 / 1511	3	16	45	47
458.sjeng	134	1 / 145.06 / 1252	0	5	11	11
462.libquantum	95	1 / 56.64 / 626	0	1	9	9
464.h264ref	523	1 / 171.42 / 5445	3	22	50	52
470.lbm	17	6 / 123.41 / 680	0	0	0	0
471.omnetpp	1406	1 / 26.9 / 611	45	69	227	270
473.astar	101	1 / 67.11 / 584	0	2	4	4
482.sphinx3	326	1 / 80 / 924	2	6	24	26
483.xalancbmk	14191	1 / 38.58 / 3809	3057	4573	4342	4887

Table 4.1 provides detailed statistics for the SPEC CPU2006. We show how many

functions (#Fns) are present in the linked program just before the merging pass, as well as the average, minimum, and maximum size of these functions, in number of instructions, at this same point in the compilation pipeline. We also report how many pairwise merge operations are performed by each one of the function merging techniques. Note that in almost all cases FMSA performs significantly more merge operations than the other techniques. There are only two cases where FMSA with exploration threshold of one finds fewer profitable merges than the state-of-the-art. This is due to our aggressive pruning of the search space with our ranking mechanism. Simply increasing the threshold, e.g. to ten, allows FMSA to merge more functions. In any case, these extra merge operations of the state-of-the-art have little effect on the overall code size reduction. The state-of-the-art is more likely to fail to merge large functions and succeed with small ones, so even when merging more functions, it does not reduce code size as much as FMSA.

MiBench: Embedded Benchmark Suite

We have shown that our technique achieves good results when applied on the SPEC CPU suite. It reduces size not only on templated C++ benchmarks, like other techniques, but also on C benchmarks where merging opportunities should be almost non-existent. Here, we further explore how FMSA handles such cases by applying it on the MiBench suite, a collection of small C programs each one composed of a small number of functions.

Figure 5.17 shows the object file reduction for the MiBench programs on the Intel platform. Our best result is for the `rijndael` benchmark, which implements the well-known AES encryption. FMSA merges the two largest functions, namely, `encrypt` and `decrypt`. Inspecting the LLVM IR for the `rijndael` benchmark, we observe that the two functions contain a total of 2494 instructions, over 70% of the code. When we merge them by sequence alignment, we create a single function with only 1445 instruction, a 42% reduction in the number of IR instructions. This translates into a 20.6% reduction in the linked object file.

Table 5.1 provides more detailed statistics for MiBench. LLVM achieves very limited results, reducing `jpeg_c` by just 0.13%, `jpeg_d` by 0.1%, and `ghostscript` by 0.02%, while having no effect on `typeset`. This happens because all the functions merged by LLVM’s identical technique are tiny functions relative to the overall size of the program. Most of these functions comprise of just a few IR instructions. For example, in the `typeset` benchmark, while it is able to merge a pair of functions, they

Table 4.2: Number and size of functions present in each MiBench benchmark just before function merging, as well as number of merge operations applied by each technique.

Benchmarks	#Fns	Min/Avg/Max Size	Identical	SOA	FMSA[t=1]	FMSA[t=10]
CRC32	4	8 / 24.75 / 39	0	0	0	0
FFT	7	7 / 49.86 / 144	0	0	0	0
adpcm_c	3	37 / 73 / 100	0	0	0	0
adpcm_d	3	37 / 73 / 100	0	0	0	0
basicmath	5	4 / 70.8 / 232	0	0	0	0
bitcount	19	4 / 22.26 / 63	0	1	3	3
blowfish_d	8	1 / 245.38 / 824	0	0	0	0
blowfish_e	8	1 / 245.38 / 824	0	0	0	0
jpeg_c	322	1 / 100.52 / 1269	2	6	8	11
jpeg_d	310	1 / 98.93 / 1269	3	6	10	10
dijkstra	6	2 / 33 / 89	0	0	0	0
ghostscript	3446	1 / 54.2 / 4218	53	53	234	250
gsm	69	1 / 97.06 / 737	0	3	8	8
ispell	84	1 / 105.51 / 1082	0	2	5	5
patricia	5	1 / 77 / 167	0	0	0	0
pgp	310	1 / 88.52 / 1845	0	1	10	10
qsort	2	11 / 50 / 89	0	0	0	0
rijndael	7	46 / 472.29 / 1247	0	0	1	1
rsynth	46	1 / 97.28 / 778	0	0	0	0
sha	7	12 / 53.29 / 150	0	0	0	0
stringsearch	10	3 / 47.9 / 99	0	0	1	1
susan	19	15 / 291.84 / 1212	0	0	1	1
typeset	362	1 / 354.47 / 12125	1	4	31	35

only have five instructions. For the same benchmark, FMSA performs several merge operations, one of them between two functions with over 500 instructions. Overall, the state-of-the-art does slightly better than LLVM’s identical technique but even in its best case it cannot reduce code size more than 0.7%.

Because these embedded benchmarks are much smaller than those in the SPEC suite, trivially similar functions are much less frequent. This is why neither the state-of-the-art nor LLVM’s identical function merging technique had any real effect on these benchmarks. Our technique can look beyond trivially similar functions which allowed it to achieve significant code size reduction on these real embedded benchmarks.




Figure 4.11: Object file size reduction for Intel on the Mibench benchmark suite. Our approach (FMSA) is the only one able to achieve a meaningful reduction on these benchmarks.




Figure 4.12: Compilation-time overhead on the Intel platform. For exhaustive exploration (not shown) the average overhead is $25\times$. Through ranking, we reduce overhead by orders of magnitude. For an exploration threshold of one, FMSA has an overhead of only 15%.

4.4.3 Compilation Overhead

Figure 5.23 shows the compilation-time overhead for all optimizations. As explained in the experimental setup, we only present results when compiling for the Intel platform. Since we cross-compile on the same machine for both targets, compilation times are very similar. We also do not include results for the oracle (exhaustive) exploration. It would be hard to visualize it in the same plot as the other configurations, since it can be up to $136\times$ slower than the baseline.

Unlike the other evaluated techniques, our optimization is a prototype implementation, not yet tuned for compilation-time. We believe that compilation-time can be further reduced with some additional engineering effort. Nevertheless, by using our ranking infrastructure to target only the single most promising equivalent function for each function we examine, we reduce compilation-time overhead by up to two orders of magnitude compared to the oracle. This brings the average compile-time overhead to only 15% compared to the baseline, while still reducing code size almost as well

as the oracle. Depending on the acceptable trade-off between compilation-time overhead and code size, the developer can change the exploration threshold to exploit more opportunities for code reduction, or to accelerate compilation.




Figure 4.13: A compilation-time breakdown isolating the percentage for each major step of the optimization ($t=1$).




Figure 4.14: Runtime overhead on the Intel platform. Performance impact is almost always statistically insignificant. For the few benchmarks affected, FMSA merges hot functions.

Figure 4.13 shows a detailed compilation-time breakdown. For each major step of the proposed optimization, we present the accumulated time spent across the whole program. To better understand the overhead of each step, we use an exploration threshold of one ($t = 1$). Because the ranking mechanism performs a quadratic operation on the number of functions, computing the similarity between all pairs of functions, it is expected that ranking would be amongst the most costly steps. However, it is interesting to notice that the sequence alignment dominates most of the compilation-time overhead, especially considering that this operation is performed only once per function, when $t = 1$. Although this operation is linear in the number of functions, the Needleman-Wunsch algorithm [44] is quadratic in the size of the functions being

aligned, both in time and space. Unsurprisingly, code generation is the third most costly step, which also includes the time to optimize the merge of the parameters. The remaining steps contribute, in total, a small percentage of all the compilation-time overhead. This analysis suggests that optimizing the sequence alignment algorithm and the ranking mechanism is key to reducing even further the overall compilation-time overhead.

4.4.4 Performance Impact

The primary goal of function merging is to reduce code size. Nevertheless, it is also important to understand its impact on the programs' execution time and the trade-offs between performance and code size reduction. Figure 4.14 shows the normalized execution time. Overall, our optimization has an average impact of about 3% on programs' runtime. For most benchmarks, there is no statistically significant difference between the baseline and the optimized binary. Only for `433.milc`, `447.dealII`, and `464.h264ref` there is a noticeable performance impact.

We take `433.milc`, which has the worst result, for discussion. For an exploration threshold value of one, we merge 58 functions. Through profiling, we discovered that a handful of them contain hot code, that is, they have basic blocks that are frequently executed. If we prevent these hot functions from merging, all performance impact is removed while still reducing code size. Specifically, our original results showed a 5.11% code size reduction and an 18% performance overhead. Avoiding merging hot functions results in effectively non-existent performance impact and a code size reduction of 2.09%. This code size reduction is still about twice as good as the state-of-the-art. As with the compilation overhead, this is a trade-off that the developer can control.

4.5 Conclusion

We introduced a novel technique, based on sequence alignment, for reducing code size by merging arbitrary functions. Our approach does not suffer from any of the major limitations of existing solutions, outperforming them by more than $2.4\times$. We also proposed a ranking-based exploration mechanism to focus the optimization on promising pairs of functions. Ranking reduces the compilation-time overhead by orders of magnitude compared to an exhaustive quadratic exploration. With this framework, our

optimization is able to reduce code size by up to 25%, with an overall average of about 6%, while introducing an average compilation-time overhead of only 15%. Coupled with profiling information, our optimization introduces no statistically significant impact on performance.

Chapter 5

Effective Function Merging in the SSA Form

While the technique proposed in Chapter 4 represents a leap forward, experiments show that FMSA fails to reduce code size in some cases where it would be intuitively expected to work. Even when handling similar functions that should be profitably merged, this algorithm may fail spectacularly, producing a merged function *larger* than the combined input functions.

Closer inspection reveals that the problem stems from the inability of this approach to handle *phi-nodes*. In SSA, *phi-nodes* merge the assignments of a single variable that arrive from different control flow paths. As such, they are closely tied to how control and data flow across basic blocks and cannot be merged without examining their control flow context. FMSA generates code directly from the aligned sequences, where control flow information has been lost, merging instructions blindly with little to no consideration for their context, so it cannot handle *phi-nodes*. It overcomes this hurdle by applying register demotion, which replaces *phi-nodes* with stack variables. This works but only by artificially increasing the size of the input functions, often by twice or more their original size, the exact opposite of what function merging tries to achieve. A final post-merging step of register promotion is supposed to reverse this code bloating but it often fails, leading to unprofitable merged functions.

Our idea is to keep the one thing that works well in FMSA, the idea of using sequence alignment on functions, and build around it a new function merging methodology that can handle directly control and data flow with no need for register demotion. Our proposed approach, SalSSA, achieves this with a new code generator for aligned functions. Instead of translating the alignment directly into a merged function, our

approach generates code from the input control-flow graphs, using the alignment only to specify pairs of matching labels and instructions. The generator then produces code top-down, starting with the control flow graph of the merged function, then populating with instructions, arguments and labels, and finally with *phi-nodes* which maintain the correct flow of data. SalSSA is carefully designed to produce correct but, still, succinct code. A final post-generator stage applies a novel optimization, *phi-node coalescing*, that eliminates superfluous phi-nodes and select instructions, reducing even further the code size.

SalSSA produces functions much smaller than those produced by FMSA. In many cases, it produces profitable merged functions where FMSA fails. On average, it reduces about twice as much code as their approach, 11.4% to 14.5% compared to 5.6% to 6.2% depending on the function merging configuration. On top of that, the compile-time overhead is much lower. Sequence alignment has a quadratic relationship with function size, while the overhead of code generation and later optimization passes is proportional to function size. By avoiding register demotion, we keep input function sequences smaller and we produce smaller functions, leading to an average compilation overhead of 5%, $3\times$ less than FMSA, and an overhead in no case more than 55%, compared to the maximum overhead of 314% for FMSA. Similarly, SalSSA uses half the amount of memory required on average by FMSA during compilation.

With this chapter, we make the following contributions:

- The first approach that fully supports the SSA form when merging functions through sequence alignment.
- A novel optimization called phi-node coalescing that reduces the number of phi-nodes and selections in merged functions.
- SalSSA achieves about twice as much code size reduction than the state of the art with significantly lower compilation time overheads.

5.1 Motivating Example

As a motivating example, consider the pair of input functions shown in Figure 5.1. While they are artificial, they highlight and isolate a problem that frequently appears in real programs, as we discuss later in Figure 5.4. These two functions have enough similarity to be profitably merged. A human expert could even replace them with the function shown in Figure 5.2, reducing the number of instructions by about 20%.




Figure 5.1: Original input functions to be merged, before register demotion. These simplified functions highlight a problem commonly seen in real programs.

However, before aligning and automatically merging them, FMSA has to apply register demotion, as shown in Figure 5.3. *Phi-nodes* are removed and memory operations are created to propagate values across basic block boundaries. The sequence alignment algorithm then identifies the matching pairs of instructions (connected green marks), keeping the rest unaligned (in red).

The problem arises when merging some of the generated memory operations. To reverse the effect of register demotion, FMSA applies register promotion on the merged code, replacing the memory operations back with *phi-nodes*. This is mandatory in FMSA in order for merged functions to be profitable, given that register demotions artificially increases the size of the functions being merged. However, in order to be promotable, a stack location must be always used directly as the immediate argument of the operations that access the location. Unfortunately, merging these instructions tend to prohibit register promotion, which results in unprofitable merge operations.

In our example, we see in Figure 5.3 that some of the mergeable memory operations use different locations. One such case is the highlighted pair of *store* instructions. To maintain the semantics of the two functions after merging, the target address of the merged *store* will have to be selected based on the function identifier, either `addr2` or `addr3`. Because the merged *store* instruction will not use the stack address directly, but instead a selected address, this prevents register promotion from eliminating these

memory operations.




Figure 5.2: Desired merged function that can be produced by an expert. An extra argument called `%fid` is used to select between the two functions. This represents a gain of about 20% in the total number of instructions.

This failure to remove temporarily inserted stack operations has knock-on effects beyond the few extra instructions left in the merged code. The additional memory accesses and the select statements controlling their target locations prohibit parts of the post-merge cleanup and later optimization passes. In our example, while the two original input functions had nine and ten instructions each, the merged function ends up with a total of 50 instructions, significantly larger than the two input functions put together.

This kind of undesired scenario is likely to happen when merging two distinct functions after register demotion simply due to the sheer number of memory operations it creates. Figure 5.4 shows the average normalized size, before and after register demotion, across all functions in each program from the SPEC CPU2006 benchmark suite. Size refers to the number of LLVM IR instructions. On average, register demotion increases function size by almost 75%, often by twice or more their original size. Even if FMSA fails to eliminate only a small portion of these extra instructions, the negative impact on the profitability of merging will be significant.

Even for cases where the merge operation is profitable, register demotion remains a problem. Demotion artificially lengthens the functions to be aligned which in turns

exacerbates the compile-time overheads associated with function merging. In our example, the combined size of the two input functions more than doubles, from 14 instructions in Figure 5.1 to 29 instructions in Figure 5.3. This increase is in line with what we have seen in SPEC CPU2006, including functions with many thousands of instructions. Regardless of whether register promotion will eventually remove the extra instructions or not, the alignment algorithm itself will have to process sequences twice as long. Since the memory usage and running time of the algorithm is quadratic in the sequence length, register demotion slows it down approximately by a factor of four. For applications with large functions after register demotion, the compile-time and memory usage overheads become prohibitive.

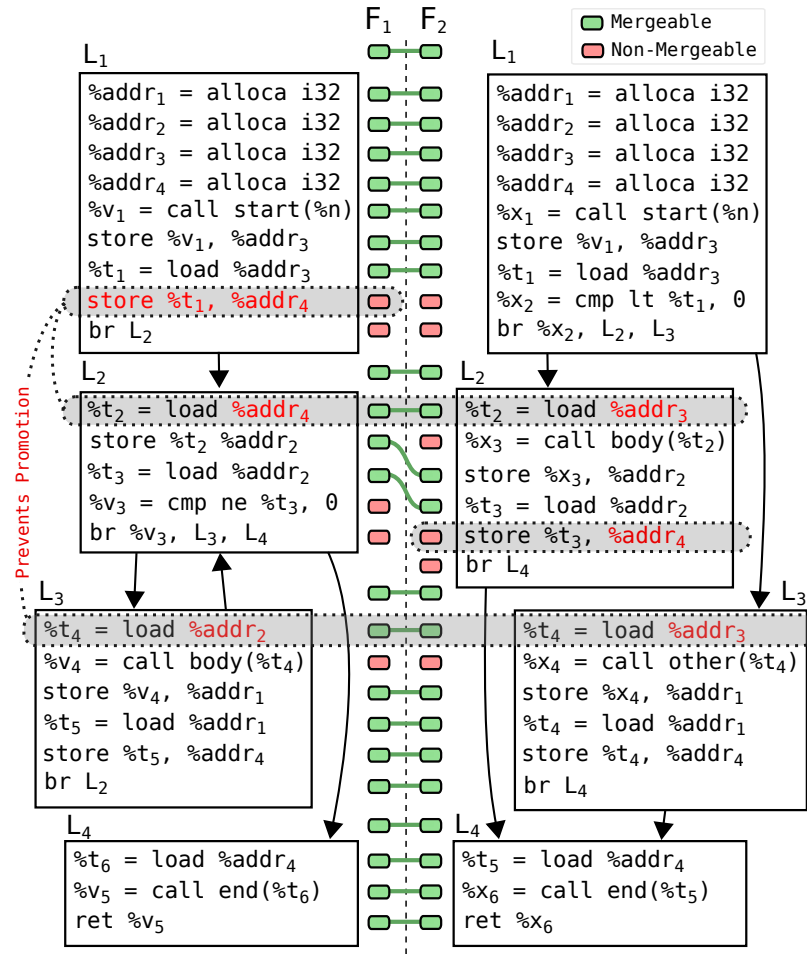


Figure 5.3: Aligned example functions after register demotion. The functions double in size after demotion, slowing down alignment. Merging some of the generated stack accesses will prevent eliminating them later through register promotion.

This shows that a new solution is needed to effectively merge functions in the SSA form. Register demotion makes function merging less profitable, even stopping

similar functions from merging altogether, and often leads to undesirable compilation overheads. In the rest of the chapter, we show that register demotion is not required for function merging and that we can directly handle *phi-nodes*, leading to more profitably merged functions.




Figure 5.4: Average normalized function size, before and after register demotion, across all functions in each program from the SPEC 2006 benchmark suite. Register demotion increases function size by almost 75% on average.

5.2 The SalSSA Approach

Properly handling *phi-nodes* requires a radical redesign in the code generator. The existing code generator produces code directly from the aligned sequence, with each instruction pair treated almost in isolation without considering any control flow context. Merging *phi-nodes* cannot work with this approach because *phi-nodes* are only understood in their control flow context.

Road map In the rest of this section, we describe SalSSA, our novel approach for merging functions through sequence alignment with full support for the SSA form. By removing the need for preprocessing the input functions and performing register demoting, our approach is able to merge functions better and faster. Instead of translating the aligned functions directly to merged code, the SalSSA follows a top-down approach centred on the CFGs of the input functions. It iterates over the input CFGs, constructing the CFG of the merged function, interweaving matching and non-matching instructions (Section 5.2.1). Afterwards, all edges and operands are resolved, including appropriately assigning the incoming values to all *phi-nodes* (Section 5.2.2). SalSSA

is designed to preserve all properties of SSA form via the standard SSA construction algorithm (Sections 5.2.3). Finally, SalSSA integrates a novel optimization with the SSA construction algorithm, called *phi-node coalescing*, producing even smaller merged functions (Section 5.2.4).




Figure 5.5: Example functions aligned without register demotion. *Phi-nodes* are excluded from alignment.

Working examples Figure 5.5 shows how the functions from our motivating example align without register demotion. Here, *phi-nodes* are not aligned, similarly to how FMSA handles *landing-pad* instructions. We will use these as working examples to describe step by step how our new code generator works in the next subsections.

5.2.1 Control-Flow Graph Generation

Our code generator starts by producing all the basic blocks of the merged function. Each original block is broken into smaller ones so that matching code is separated from non-matching code and matching instructions and labels are placed into their own basic blocks. Having one block per matching instruction or label makes it easier to handle control flow and preserve the ordering of instructions from the original functions by chaining these basic blocks as needed.

Blocks with instructions that come originally from the same basic block (of either input function) are chained in their original order with branches. We use either un-

conditional branches or conditional branches on the function identifier depending on whether control flow out of this code is different for the two input functions. Because we have one basic block per pair of matching instructions/labels, this tends to generate some artificial branches, most of them are unconditional, but can be simplified in later stages.




Figure 5.6: Merged CFG produced by SalSSA. Code corresponding to a single input basic block may be transformed into a chain of blocks, separating matching and non-matching code. The generator inserts conditional and unconditional branches to maintain the same order of instructions from the input basic block. Operands and edges highlighted in blue will be resolved by the operand assignment described in Section 5.2.2.

Figure 5.6 shows the generated CFG. At this point, the only instructions that actually have their operands assigned are the branches inserted to chain instructions originating from the same input basic block. These branches have no corresponding in-

struction in the input functions. All other operands and edges, depicted in blue in Figure 5.6, will be resolved later, during operand assignment.

5.2.1.1 Phi-Node Generation

Our code generator treats *phi-nodes* differently from other instructions. For all alignment and code generation purposes, SalSSA treats *phi-nodes* as attached to their basic block's label; that is, they are aligned with their labels and are copied to the merged function with their labels. So, when creating a basic block for a label, we also generate the *phi-nodes* associated with it. For a pair of matching labels, we copy all *phi-nodes* associated with both labels. We have decided for this approach where *phi-nodes* are tied to labels because *phi-nodes* describe primarily how data flows into its corresponding basic block. Figure 5.6 shows an example where *phi-nodes* are present in basic blocks with both matching or non-matching labels. The *phi-node* x_5 is simply copied into the merged basic block labeled L_6 .

Unlike other instructions, we do not merge *phi-nodes* through sequence alignment. Instead, identical *phi-nodes* are merged during the simplification process using existing optimizations from LLVM.

5.2.1.2 Value Tracking

While generating the basic blocks and instructions for the merged function, SalSSA keeps track of two mappings that will be needed during operand assignment. The first one, called *value mapping*, is responsible for mapping labels and instructions from the input functions into their corresponding ones in the merged function. This is essential for correctly mapping the operand values. The second one, called *block mapping*, is a mapping of the basic blocks in the opposite direction, as shown by the light gray labels in Figure 5.6. It maps basic blocks in the merged function to a basic block in each input functions, whenever there is a corresponding one. This *block mapping* will be needed to map control flow when assigning the incoming values of *phi-nodes* (see Section 5.2.2.3).

5.2.2 Operand Assignment

Once all instructions and basic blocks have been created, we perform operand assignment in two phases. First, we assign all label operands, essentially resolving the remaining edges in the control flow graph (dashed blue edges in Figure 5.6). With the

control flow graph complete, we can then create a dominator tree to help us assign the remaining operands while also properly handling instruction domination.

```

%m2 = call body(%v2|%m1)
      |||
%s = select %fid==1, %v2, %m1
%m2 = call body(%s)

```

Figure 5.7: Operand selection for the `call` instruction in L_4 from Figure 5.6. Mismatching operands chosen with a `select` instruction on the function identifier.

Whenever the corresponding operands of merged instructions are different, we need a way to select the correct operand based on the function identifier. Section 5.2.2.1 describes how we perform label selection. In all other cases, we simply use a `select` instruction, as shown in Figure 5.7.

```

%y = add %m|%b1, %a2|%m
      |||
%s = select %fid==1, %a2, %b1
%y = add %m, %s

```




Figure 5.8: Optimizing operand assignment for commutative instructions. Example of a merged `add` instruction that can have its operands reordered to allow merging the two uses of `%m`, avoiding a `select` instruction.

When assigning operands to commutative instructions, we also perform operand reordering to maximize the number of matching operands and reduce the need for `select` instructions. Figure 5.8 shows an example of a commutative instruction where an operand selection can be avoided by reordering operands. This property of commutative operations has been exploited before by other optimizations [47, 48, 49].

5.2.2.1 Label Selection

In LLVM, labels are used exclusively to represent control flow. More specifically, label operands are used by terminator instructions, where they specify the destination basic block of a control flow transfer, or to represent incoming control flow in a *phi-node* instruction.

Whenever assigning the operands of a merged terminator instruction, if there is a label mismatch between the two input functions, we need a way to select between

the two labels depending on the executed function. We do so by creating a new basic block with a conditional branch on the function identifier to each one of the mapped labels. Then we use the new block's label as the operand of the merged terminator instruction. Figure 5.9 illustrates a CFG that handles label selection for a merged terminator instruction.




Figure 5.9: Label selection for mismatched terminator instruction operands L_{f1} and L_{f2} corresponding to labels of two different basic blocks. We handle control flow in a new basic block, L_{sel} with a conditional branch on the function identifier targeting the two labels. We use the label of the new block as the merged terminator operand.

Figure 5.10 shows a special case where we can also perform operand reordering on conditional branches that follow a specific pattern. When merging two conditional branches with matching label operands, except for their order, instead of creating two label selections, we can simply apply an *xor* operation on the condition and the function identifier, swaping the label operands for the true-value of the function identifier. As shown in Figure 5.10b, the *xor* operation flips the value of the condition for the true-value of the function identifier, preserving the semantic of the conditional branch. This optimization adds the cost of one *xor* operation to avoid the cost of two label selections, which are implemented with branch instructions as shown in Figure 5.9.

5.2.2.2 Landing Blocks

Most modern compilers, including GCC and LLVM, implement the zero-cost Itanium ABI for exception handling [14], which is known as the *landing-pad* model. This model has two main components: (1) invoke instructions that have two successors, one that continues when the call succeeds as per normal, and another, usually called the *landing pad*, in case the call raises an exception, either by a throw or the unwinding of a throw; (2) landing-pad instructions that encode which action is taken when an exception is raised. A landing pad must be the immediate successor of an invoke

$\text{br } \%C, L_1, L_2, L_2, L_1$
 $\text{\%xc} = \text{xor } \%C, (\%fid==1)$
 $\text{br } \%xc, L_2, L_1$

(a) Rule for conditional branches with swapped label operands.

%C	%fid==1	%xc
0	0	0
0	1	1
1	0	1
1	1	0

(b) The truth table of the xor operation.

Figure 5.10: Optimizing label assignment for conditional branches. Example of a merged `br` instruction that can have its label operands reordered, trading two label selections by one xor operation.





Figure 5.11: Landing blocks are added after operand assignment and are assigned to invoke instructions as operands.

instruction in its unwinding path. The code generator must ensure that this model is preserved.

Our new code generator delays the creation of landing-pad instruction until the phase of operand assignment. Once we have concluded the remapping of all label operands of an `invoke` instruction, regardless of whether they are merged or non-merged code, we create an intermediate basic block with the appropriate landing-pad instruction. Then we assign the label of this landing block as the operand of the `invoke` instruction, as shown in Figure 5.11.

5.2.2.3 Phi-Node's Incoming Values

There are two distinct cases for *phi-nodes*: being associated with a matching or with a non-matching label. In both cases, *phi-nodes* are only copied from their input functions and they are not merged. So each *phi-node* in the merged function should capture the incoming flows present in the corresponding *phi-node* of their input function. For



(a) Example where the dominance property is violated. (b) The dominance property is restored by placing phi-nodes where needed.

Figure 5.12: Example of how SalSSA uses the standard SSA construction algorithm to guarantee the dominance property of the SSA form.

matching labels, each *phi-node* in the merged function will have additional incoming flows specific to the *other* input function but these flows should have undefined values.

To assign the incoming values of a *phi-node*, SalSSA iterates over all predecessors of its parent basic block and uses the *block mapping* to discover each predecessor's corresponding basic block in the input function. If such a basic block is found, then SalSSA obtains the incoming value associated with that predecessor from the *value mapping*. Otherwise, an undefined value, which by construction should never be actually used, is associated with that predecessor.

5.2.3 Preserving the Dominance Property

The code transformation process described so far could violate the *dominance property* of the SSA form. This property states that each use of a value must be dominated by its definition. For example, an instruction (or basic block) dominates another if and only if every path from the entry of the function to the latter goes through the former. Figure 5.12a gives one example extracted from Figure 5.6 where the dominance property is violated during code transformation. For this example, the dominance property is violated because $\%v_2$ is defined in block L_{12} and used in block L_4 , but the former does not dominate the latter since there is an alternative path through L_{21} .

SalSSA is designed to preserve the dominance property to conform with the SSA form. It achieves this using a two-step approach. It first adds a *pseudo-definition* at the entry block of the function where names are defined and initialized with an *undefined* value. This guarantees that every register name will be defined on basic

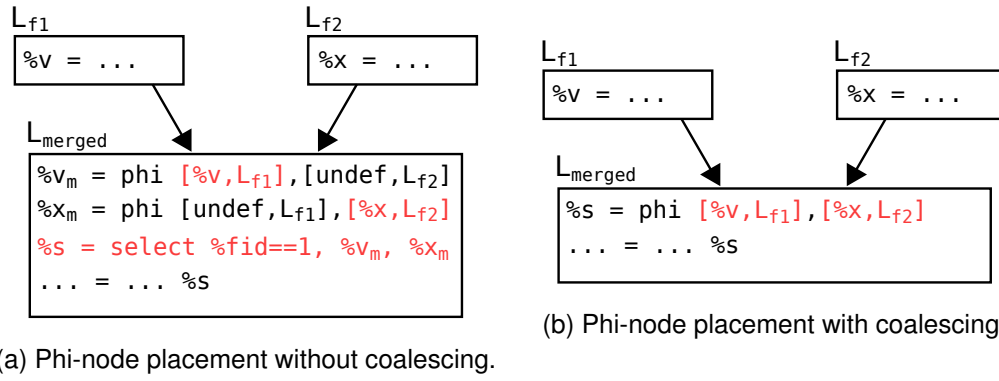


Figure 5.13: Phi-node coalescing reduces the number of phi-nodes and selections.


blocks from both functions. Then, SalSSA applies the standard SSA construction algorithm [11, 12], which guarantees both the dominance and the single-reaching definition properties of the SSA form. We note that our implementation uses the standard SSA construction algorithm provided by LLVM for register promotion. This algorithm guarantees that names have a single definition by placing extra phi-nodes where needed so that instructions can be renamed appropriately. Figure 5.12b shows how the property violation in Figure 5.12a can be corrected using this strategy.

5.2.4 Phi-Node Coalescing

The approach described in Section 5.2.3 guarantees the correctness of the SSA form but generates extra phi-nodes and registers which increase register pressure and might lead to more *spill code*. In this section, we describe a novel optimization technique, *phi-node coalescing*, that SalSSA uses to lower register pressure.

Figure 5.13 illustrates such an optimization opportunity. SalSSA is merging an instruction with different arguments, so it needs to select the right one based on the function identifier. The two arguments though, v and x , have *disjoint definitions*, i.e. they have non-merged definitions from different input functions. Using the standard SSA construction algorithm would result in the sub-optimal code shown in Figure 5.13a. This code inserts two trivial phi-nodes to select, again, v or x based on the executed function. SalSSA optimizes this code by coalescing both phi-nodes into a single one and removing the selection statement. As shown in Figure 5.13b, the optimized version has a smaller number of instructions and phi-nodes.

This transformation is valid because a value definition that is exclusive to a function will never be used when executing the other function. Figure 5.14 shows another



(a) Phi-node placement without coalescing.

(b) Phi-node placement with coalescing.

Figure 5.14: Reducing the number of phi-nodes by coalescing disjoint definitions with no user instructions in common.

example illustrating that even disjoint definitions that have no user instructions in common can be coalesced, reducing the number of phi-nodes.

Since SalSSA is aware of which basic blocks are exclusive to each function, it can choose a pair of disjoint definitions for coalescing. Given a pair of disjoint definitions, SalSSA assigns the same name for both of them before applying the SSA reconstruction. SalSSA coalesces the set of definitions that violate the dominance property. Two definitions can be paired for coalescing if they are disjoint and have the same type. The optimization pairs disjoint definitions that maximize their live range overlap since the goal is to avoid having register names live longer than they should, reducing register pressure.

Formally, the heuristic implemented in our phi-node coalescing can be described as follows. Given a set $S_1 \times S_2$ of disjoint definitions that violate the dominance property, the optimization chooses pairs $(d_1, d_2) \in S_1 \times S_2$ that maximize the intersection $UB(d_1) \cap UB(d_2)$, where $UB(d)$ is the set $\{Block(u) : u \in Users(d)\}$.

Phi-node coalescing allows SalSSA to produce smaller merged functions and reduce code size. Consequently, it also enables more functions to be profitably merged.

5.3 Evaluation

In this section, we compare SalSSA against the function merging by sequence alignment, FMSA, introduced in Chapter 4. We first present the code size reduction on the




Figure 5.15: Compilation pipeline used for the evaluation. Both SalSSA and FMSA are applied in LTO mode.

final object file. We then evaluate the compilation overhead and impact on program performance.

5.3.1 Experimental Setup

Most of our experiments directly compare SalSSA against FMSA [49]. We use the same compilation pipeline as our prior work [49], depicted in Figure 5.15. Both function merging optimizations are implemented in LLVM version 11.

Our approach uses the same fingerprint-based ranking mechanism as FMSA to decide which functions to attempt to merge. This strategy uses a configurable *exploration threshold*, t , to control how many different functions to attempt to merge with each function before selecting the most profitable merge or give up. A larger exploration threshold (t) is likely to lead to better code size reduction, but comes at the cost of longer compile time. Like FMSA, we also use three different exploration thresholds where $t = \{1, 5, 10\}$.

We evaluated SalSSA and FMSA on all C/C++ benchmarks of the SPEC CPU benchmark suite [55], both the 2006 and 2017 versions, targeting the Intel x86 architecture, and on the MiBench embedded benchmark suite targeting the ARM Thumb architecture. We run all experiments on a dedicated server with a quad-core Intel Xeon CPU E5-2650, 64 GiB of RAM, running Ubuntu 18.04.3 LTS. To minimize the effect of measurement noise, compilation and runtime overhead experiments were repeated 5 times.

5.3.2 Evaluation on SPEC CPU

Figure 5.16 reports the code size reduction on linked objects over the LLVM link-time optimizer (LTO). SalSSA significantly improves FMSA. With the lowest exploration threshold, SalSSA on average reduces the compiled code size by 9.3% and 7.9% on




Figure 5.16: Linked object size reduction over LLVM LTO when performing function merging with SalSSA or FMSA on SPEC CPU 2006 (a) and 2007 (b). Each approach was evaluated using three different exploration thresholds. On SPEC CPU2006, SalSSA reduces code size by 9.3% to 9.7% on average, almost twice as much as FMSA. On SPEC CPU2017, SalSSA reduces code size by 7.9% to 9.2% on average, more than twice as much as FMSA.

SPEC CPU2006 and CPU2017 respectively. These translate to nearly twice or above as much as FMSA, which achieves a 3.8% and 4.1% reduction on SPEC CPU2006 and CPU2017 respectively. The highest reductions are seen for 447.dealII and 510.parset_r, over 40% reduction. They are mainly due to the heavy use of template functions which leads to multiple similar functions. Other C++ programs display similar behavior, where SalSSA also achieves good code size reduction. SalSSA also gives remarkable code size reduction in many C programs, such as 456.hmmer, 462.libquantum, and 482.sphinx3.

SalSSA outperforms FMSA for multiple benchmarks. The more pronounced cases are for 444.namd, 456.hmmer, 462.libquantum, 447.dealII, and 482.sphinx3 from SPEC CPU2006, as well as 508.namd_r, 619.lbm_s, 644.leela_s and 657.xz_s from SPEC CPU2017. These benchmarks were heavily affected by register demotion, as illustrated in Figure 5.4 for SPEC CPU2006. Similar to our motivating example in Section 6.1, when two non-identical functions have stack operations for nearly half of

their instructions, misalignments become likely; these misalignments prohibit eliminating the merged stack operation through register promotion. This issue reduces the profit gained by FMSA. In some cases, like `619.lbm.s` and `625.x264.s`, the profitability cost model can fail, resulting in sufficient false positives to cause code bloating. We will discuss this further in the next section.

5.3.3 Evaluation on MiBench

To evaluate the effectiveness of SalSSA on embedded systems, we apply it to the MiBench embedded benchmark suite on the ARM Thumb architecture. We note that by having function merging implemented at the IR level, our approach can be equally applied to any target architecture supported by the compiler.

The MiBench suite is a collection of short C programs, each one composed of a small number of functions. When optimizing programs with a small number of functions, function merging optimizations will have fewer opportunities to find pairs of profitably merged functions. For example, the `qsort` program in MiBench has only two functions; as a result, neither FMSA nor SalSSA is able to merge them. As shown in Table 5.1, the same happens for other programs in the MiBench suite.

Figure 5.17 shows that SalSSA improves significantly over FMSA, achieving a geo-mean reduction of 1.4% to 1.6%, about twice as much as FMSA. This improvement comes from SalSSA’s capability of generating better-merged functions, which leads to a larger number of profitable merge operations, as confirmed by Table 5.1.

Because FMSA requires register demotion to be applied to all functions before it can even attempt to merge them, FMSA ends up changing all functions even if no profitable merge operation is found. Figure 5.17 shows the effect of this preprocessing phase (denoted as *FMSA Residue*), which is obtained by running FMSA but not committing any merge operation. This FMSA Residue is the reason why FMSA sometimes has a non-zero code-size reduction (e.g., `adpcm_c`, `FFT`, `patricia`) despite not merging any functions. Since FMSA Residue might have an impact on the heuristics of later optimizations and code generation, its impact is almost random, sometimes being positive or negative on code-size. The impact of FMSA Residue is more noticeable in small programs, such as those found in MiBench, while in SPEC2006 it increases code size by only 0.02%, on average. To fix the issue highlighted by FMSA Residue, we would need to add an extra bookkeeping step of cloning all original functions so that we can rollback if they are not profitably merged. Fixing that would only increase

Table 5.1: Number and size of functions present in each MiBench benchmark just before function merging, as well as number of merge operations applied by each technique.

Benchmarks	#Fns	Min/Avg/Max Size	FMSA[t=1]	SalSSA[t=1]
CRC32	4	8/23.75/37	0	0
FFT	7	6/45.43/131	0	0
adpcm_c	3	35/68.33/93	0	0
adpcm_d	3	35/68.33/93	0	0
basicmath	5	4/60/204	0	0
bitcount	19	4/20.58/56	3	3
blowfish_d	8	1/231.38/790	0	1
blowfish_e	8	1/231.38/790	0	1
cjpeg	322	1/92.76/1198	7	26
dijkstra	6	2/31.5/83	0	0
djpeg	310	1/91.31/1198	10	28
ghostscript	3452	1/50.36/3749	211	327
gsm	69	1/92.42/696	6	9
ispell	84	1/97.08/1004	3	8
patricia	5	1/73.6/160	0	0
pgp	310	1/80.39/1706	8	19
qsort	2	11/45.5/80	0	0
rijndael	7	45/444.14/1182	1	1
rsynth	47	1/83.89/716	1	2
sha	7	12/49.71/147	0	1
stringsearch	10	3/41/81	1	1
susan	19	15/275.21/1153	1	2
typeset	362	1/327.61/11744	27	53

unnecessarily the optimization complexity, but SalSSA offers a better solution where only merged functions are affected.

An interesting case is observed with both `cjpeg` and `djpeg`. Although SalSSA, with exploration threshold $t = 1$, increases code size, it is merging a superset of the pairs of functions merged by FMSA with $t = 1$. If we limit SalSSA to merge exactly the same pairs merged by FMSA, it ends up with about the same or slightly better results than FMSA. This suggests that the marginal code-size increase observed with SalSSA is a result of false positives from the profitability cost model, i.e., it allows unprofitable merge operations to be committed. Since `cjpeg` and `djpeg` share most of their code base, we can indeed confirm that a subset of the pairs of functions merged by SalSSA, for both benchmarks, should have been classified as unprofitable as merging them increases the code size. However, with higher exploration thresholds, namely, $t = 5$ and $t = 10$, SalSSA surpasses FMSA in code-size reduction, although it still




Figure 5.17: The percentual reduction in size of the linked object files, targeting the ARM architecture. We evaluate SalSSA or FMSA over the LLVM LTO on the MiBench embedded benchmark suite. Each approach was evaluated using three different exploration thresholds. SalSSA achieves a geo-mean reduction of 1.4% to 1.6%, about twice as much as FMSA.



Figure 5.18: A breakdown of SalSSA[$t = 1$] on the `djpeg` benchmark. The actual contribution to the final code size for each merge operation deemed profitable by the cost model.

includes all pairs merged with the exploration threshold $t = 1$.

Figure 5.18 shows a breakdown for each merge operation performed by SalSSA, with exploration threshold $t = 1$, on the `djpeg` benchmark. We measured the impact of each merge operation, in isolation, to the size of the final object file. Although each one of these merge operations have a very small contribution to the final code size, the profitability cost model failed enough to result in an overall code increase of about 0.3%.

Both SalSSA and FMSA use the same profitability cost model. The limitations observed on `cjpeg` and `djpeg` also appear in SPEC2017 with FMSA. This stems from the fact that several transformations will still be applied to the code during late optimizations and the back end, and these changes are not captured by the profitability cost model.




Figure 5.19: Evaluation of the impact of phi-node coalescing on the size of the final object file. SalSSA-NoPC, which includes phi-node coalescing, has a measurable benefit over the alternative without phi-node coalescing (SalSSA-NoPC). When enabled, phi-node coalescing achieves up to 7% of code size reduction on top of SalSSA-NoPC.

5.3.4 Further Analysis

We also provide a breakdown showing the impact of phi-node coalescing on code size. Figure 5.19 shows the impact of our phi-node coalescing optimization technique (see Section 5.2.4). This diagram compares SalSSA to a variant without phi-node coalescing (SalSSA-NoPC) and FMSA. On average, this technique gives an additional 1.2% on code size reduction. For `444.namd`, it enables an extra 7% reduction on the code size, demonstrating the great advantage of the technique.

Figure 5.20 provides further insight into the gains of SalSSA. The figure shows the total number of profitable merging attempts for the lowest exploration threshold. While FMSA has only 9,271 profitable merge operations, SalSSA has 12,224, an increase of 31% on the number of profitable merges. Much of the improvement we observe in code size reduction comes from producing profitable merged functions where FMSA fails to gain any profit, not just from increasing the profit.

5.3.5 Memory Usage

Because the sequence alignment algorithm [44, 49] (used by FMSA and SalSSA) has a quadratic space complexity over the length of the sequences, the difference in the size of the functions caused by register demotion translates directly to the differences in memory usage.

Figure 5.21 shows the peak memory usage across the SPEC CPU2006 suite. To iso-




Figure 5.20: Total number of profitable merge attempts for SalSSA and FMSA on 19 SPEC CPU2006 benchmarks. For both cases, we used the lowest exploration threshold ($t=1$). SalSSA achieves 31% more profitable merge operations.




Figure 5.21: Peak memory usage during compilation time on the SPEC CPU2006 benchmark. On average, SalSSA requires less than half the memory used by FMSA.

late the impact of other compilation passes, we measure the memory usage only when running the function merging optimization. As expected, avoiding register demotion has the added benefit of lowering the memory footprint of the compilation pass. On average, SalSSA uses half the amount of memory required by FMSA. The improvements on memory usage shown in this Figure 5.21 directly reflects the difference shown in Figure 5.4.

Both FMSA and SalSSA starts merging from the largest to the smallest functions. For the 403.gcc benchmark, the first pair of functions considered for a merging is the pair `recog_16` and `recog_26` that originally contains 20,688 and 16,043 instructions, respectively, but after register demotion grow to 36,508 and 28,899. This pair




Figure 5.22: Speedup over the accumulated time spent on both sequence alignment and code generation. SalSSA produces significantly less overhead than the FMSA.

of extremely large functions is responsible for the peak in memory usage when optimizing this benchmark. FMSA uses a total of 6.5 GB of memory while SalSSA is able to reduce it down to 2.4 GB. A total of $2.7\times$ reduction on peak memory usage. Although this is the most critical benchmark in terms of absolute numbers, a similar trend appears in most of the other benchmarks. By reducing the memory overhead of compilation, SalSSA thus can target a larger codebase over FMSA.

5.3.6 Compilation Time Overhead

Figure 5.23 shows the normalized compile-time for *the entire* compilation process on SPEC CPU2006. The min-max bar in the diagram gives the 95% confidence interval across different compile-time measurements of a benchmark. SalSSA incurs modest compile-time overhead with an average 5% increase in the compile-time when using the lowest exploration threshold ($t = 1$). This represents a 3x reduction in the compile-time overhead compared to the 14% overhead from FMSA with the same exploration threshold. When using the largest exploration threshold ($t = 10$), we observe a 3.7x reduction in the compile-time overhead. The improvement is due to not only less time spent performing the optimization itself but also less work for the remaining compilation process since we reduce the size of the produced code. We also observe similar overhead improvement on SPEC CPU2017.

Figure 5.22 shows the speedups obtained by SalSSA for the sequence alignment and the code generator. These two stages of function merging benefit most from our techniques. As suggested earlier, both stages are accelerated because the compiler




Figure 5.23: End-to-end compile-time for SalSSA and FMSA for three different exploration thresholds and 19 different SPEC CPU2006 benchmark. Compile-time is normalized to that of the baseline with no function merging. SalSSA reduces the overhead of function merging by $3\times$ to $3.7\times$ on average.




Figure 5.24: Comparison between the runtime impact from FMSA and SalSSA. Our approach increases the runtime overhead because it merges more functions. For most benchmarks, the overhead is small. For the rest, profiling-directed merging would eliminate the overhead.

has shorter sequences to operate on under SalSSA over FMSA. The results given in Figure 5.22 and Figure 5.4 follow a very similar trend. These confirm our intuition described earlier in Section 6.1.

Since the sequence alignment algorithm is also quadratic in time over the length of the sequences, we get a quadratic speedup by avoiding register demotion with SalSSA. Code generation is linear on the size of the functions resulting in proportional speedups in compile-time. For a couple of cases, however, the pressure put on the clean-up phase can negate those gains.

5.3.7 Performance Overhead

The primary goal of function merging is to reduce code size. Nevertheless, it is also important to keep the impact on the program runtime as low as possible. Figure 5.24 shows the normalized execution time, where the min-max bar shows the 95% confidence interval across different runs. Overall, SalSSA has an average overhead of about 4% on programs' runtime. For most benchmarks, there is no statistically significant difference between the baseline and the optimized binary. For the rest, profiling information could be used to avoid adding overhead when mergeable code is in the most frequently executed code path.

5.4 Conclusion

We have presented SalSSA, a novel compiler-based function merging technique with full support for the SSA form. Unlike FMSA, which has to apply register demotion to eliminate the commonly used *phi-nodes* in SSA, SalSSA directly processes *phi-nodes* using a more powerful code generator. As a result, SalSSA avoids the code bloating problem introduced by register demotion and increases the chances of generating profitable merged functions. We have implemented SalSSA in LLVM and evaluated it on the SPEC CPU2006 and CPU2017 benchmark suites. SalSSA delivers on average 9.5% code reduction for the lowest exploration threshold. Compared to FMSA, SalSSA achieves $2\times$ more reduction on binary size with $3\times$ less compile-time overhead and less than half the amount of memory required by it.

Chapter 6

Function Merging for Free

The technique proposed in Chapter 5, SalSSA, can lead to 40% slower compilation, taking up to 32 GB of memory for temporary data when compiling a modestly-sized program. Such a resource requirement is beyond what is typically available to a developer and thus unsuitable for optimizing real-life programs.

These inefficiencies stem directly from SalSSA’s core innovation, i.e., the sequence alignment algorithm used to identify mergeable instructions in a pair of input functions. The alignment algorithm has quadratic time and space complexity, so applying it on whole functions with thousands of instructions results in unacceptable overheads. This severely limits the applicability of function merging on relatively large programs. To make function merging scalable and practical, we need to find ways to significantly reduce the memory and compilation overhead. In this chapter, we offer such capabilities.

We present HyFM, a novel function merging technique that addresses the performance inefficiencies of SalSSA. Our main insight is that most of the code reduction of SalSSA comes from matching highly similar basic blocks. Even though it is able to align arbitrary subsequences spanning basic block boundaries, profitable alignments usually contain instructions from one block matched to instructions from a single other block. We show that an approach which quickly identifies similar basic blocks and then aligns their short instruction sequences achieves a similar code reduction for a much lower overhead.

More specifically, our solution is three fold:

- We align the input functions on a per basic block manner. First, we pair similar basic blocks by minimizing the distance between their fingerprints. Then, we only align the instructions within each pair of basic block. Even with a quadratic alignment algorithm, basic blocks are usually much shorter than functions, trans-

lating into a much faster alignment.

- We propose a linear pairwise alignment as an alternative to the quadratic one. For highly similar basic blocks, it achieves similar results but has negligible time and space overheads.
- We estimate the profitability of the aligned basic blocks before actually generating their merged code. If unprofitable, we ignore them, improving the overall profitability of the whole merged function and simplifying code generation. If all paired blocks in a pair of functions are unprofitable, we skip merging the function pair altogether, speeding up the optimisation process compared to SalSSA.

Experimental results on SPEC CPU 2006 and 2017 show that HyFM runs over $4.5\times$ faster than SalSSA. Compared to a baseline without function merging, HyFM reduces end-to-end compilation time by up to 18% and 2.1% on average. HyFM also has orders of magnitude lower peak memory usage, using up to 48 MB or 5.6 MB, depending on the variant used, while SalSSA requires 32 GB in the worst case. We achieve all these compilation-time benefits without degrading its ability to reduce code size.

6.1 Background and Motivation

In this section, we first provide an overview of the working mechanism of SalSSA, described in Chapter 5. We highlight the main drawbacks of SalSSA in terms of compile time and memory footprint. We then outline how we can address these drawbacks without compromising on code size reduction.

6.1.1 Function Merging via Sequence Alignment

Existing function merging techniques consist of three major stages: choosing which functions to merge, producing the merged function, and estimating the merging profitability.

In order to pair similar functions for merging, SalSSA employs a ranking strategy based on the similarity of the *fingerprints* of the functions. A fingerprint summarizes the content of a function as a fixed-size vector of the frequency of each LLVM-IR opcode. The representation allows the compiler to compare functions using a simple distance metric, such as the Manhattan distance. For a given reference function, all

other functions are ranked based on their distance and the closest function is chosen for merging.

Merging two functions requires identifying similar code segments in the two functions that can be profitably merged. The main innovation of SalSSA [50] and FMSA [49] is the use of a sequence alignment algorithm, called the Needleman-Wunsch algorithm, for identifying similar code segments. This allows them to merge arbitrary pairs of functions. First, they transform each function into a linear sequence of labels and instructions. Then, the alignment algorithm is applied on the sequences of the whole input functions. The resulting alignment is used to generate the merged function. Once the merged function has been generated, they apply an SSA reconstruction algorithm. For a final clean up, they simplify the merged function by removing redundant instructions introduced by function merging.

Finally, a profitability analysis estimates the benefit of replacing the original pair of functions with the simplified merged function. If unprofitable, the merged function is simply thrown away. Otherwise, they delete the original functions, redirecting the calls to the merged function.

6.1.2 Limitations of SalSSA

After further investigation, we observed that SalSSA was unable to optimize `602.gcc.s`, from SPEC 2017, due to an out of memory crash. Our machine with 16 GB of memory was not enough to handle SalSSA. We succeeded only after migrating to a 64 GB machine which could fit the 32 GB of temporary data produced by function merging. We realized that this is due to the quadratic algorithm used for aligning the two functions selected for merging. Because this algorithm is applied on the linearized sequences of the whole input functions, SalSSA incurs a high memory footprint when merging even medium sized functions. For larger ones, it is impossible to apply it on most workstations or even many servers, making SalSSA impractical for use in production.

For the same reason, alignment brings the compilation process to a crawl for large functions. Figure 6.1 shows the running time breakdown for the different stages of the function merging pass in the LLVM-based SalSSA implementation for two SPEC CPU2017 benchmarks. Sequence alignment dominates the running time of function merging, representing up to 83% of its overall running time. Sequence alignment alone takes 25 seconds and 4.2 minutes on `638.imagick.s` and `602.gcc.s`, respectively. The alignment stage also causes the peak in memory usage for these two programs,

4.5 GB for `638.imagick_s` and 32 GB for `602.gcc_s`. This is not surprising, as the Needleman-Wunsch algorithm has a quadratic complexity in both time and memory usage. Because this algorithm is applied on linearized sequences of the whole input functions, programs containing large functions, such as the ones in our example, are heavily affected.




Figure 6.1: Breakdown of the relative runtime for the different stages from SalSSA. Alignment takes 25 seconds and 4.2 minutes on `638.imagick_s` and `602.gcc_s`, respectively.

Most of the rest of the running time of function merging is associated with producing merged functions from the aligned sequences. This includes the time spent on the code generation stage (Code-Gen), SSA reconstruction (SSA Fix), and code simplification (Simplification). These stages account for 18.7% of the SalSSA’s running time on `638.imagick_s` and 11.6% on `602.gcc_s`. However, for other programs, these stages may represent the vast majority of SalSSA’s running time (see Section 6.3.3).

This breakdown includes the cost for producing both *profitable* and *unprofitable* merged functions. In fact, most of it is wasted on merged functions that will be rejected by the profitability analysis. These costs are pronounced because unprofitably merged functions have no limit on their size or complexity, often adding a significant pressure on the SSA reconstruction and simplification stages. This effect is tied to the alignment strategy, since a good alignment is needed for producing profitably merged functions. As we discuss in Section 6.1.3, a better approach would include a finer grain profitability analysis that would allow us to bail out from merging complex and unprofitable code as early as possible.

6.1.3 When Less is More

We observe that most of the benefit of function merging often comes from merging highly similar, but not necessarily identical, basic blocks. Figure 6.2 shows one such example extracted from the `483.xalancbmk` benchmark found in SPEC CPU2006. This example shows the two input functions annotated with the alignment produced by SalSSA. Merging these two functions contributes to a reduction of 33 bytes in the final object file.

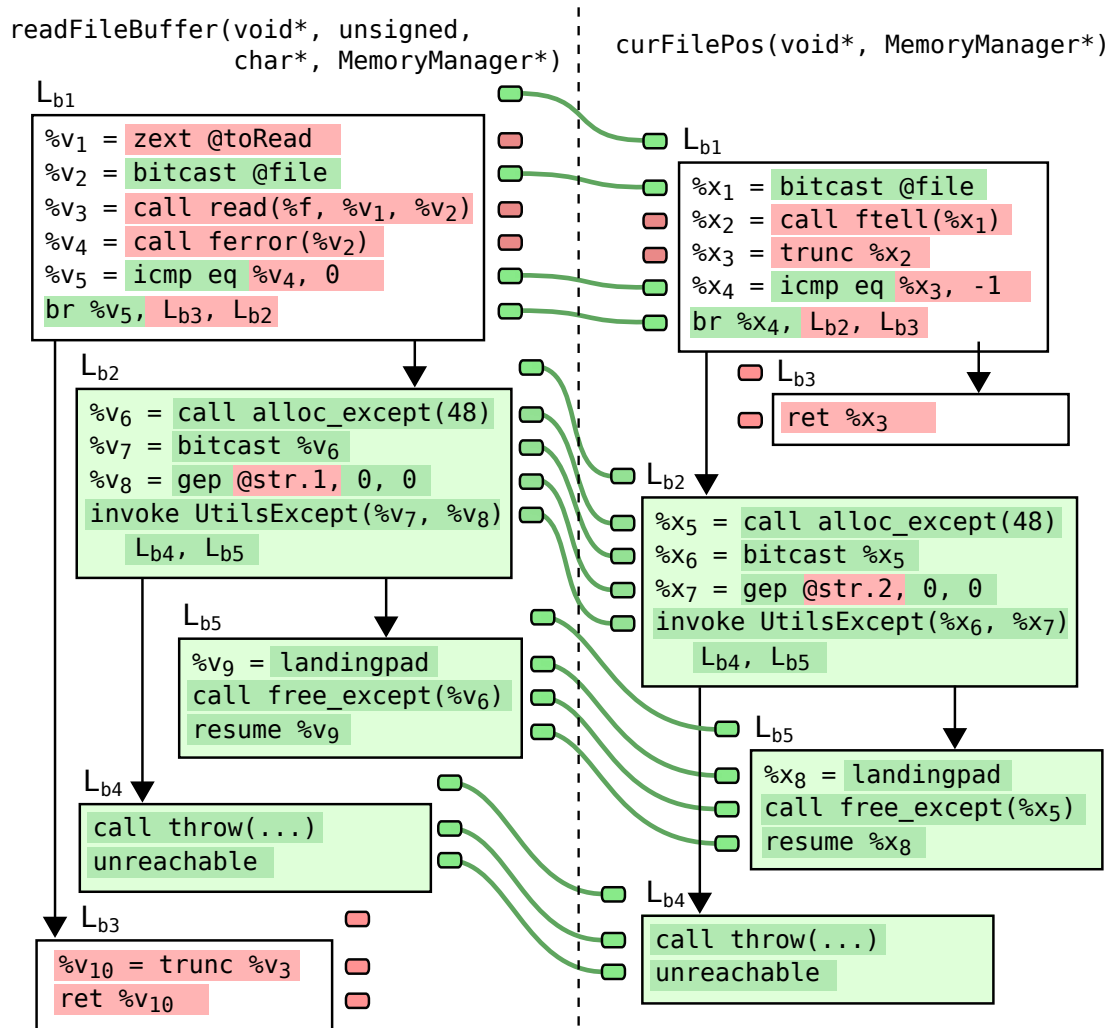


Figure 6.2: Example extracted from `483.xalancbmk` in SPEC CPU2006. Instructions marked green have been aligned through sequence alignment with an instruction from the other function. SalSSA would attempt merging all matched instructions but only the ones in fully aligned basic blocks would be profitable.

While this approach is flexible enough to identify very complex alignments, what it actually produces is three aligned pairs of basic blocks and a few aligned instructions

in the entry blocks. More importantly, these entry block instructions offer nothing in terms of code size reduction. The gains of merging them are negated by the extra branches and operand selections needed to preserve the program’s semantics. Since SalSSA analyzes the profitability of the final merged function as a whole, this unprofitable sequence of instructions will be merged because of the three highly profitable basic blocks. For the same reason, we may have profitable areas of code rejected because the rest of the merged function is unprofitable.

This example shows us that we could achieve similar code size reduction by breaking the problem of aligning functions into two simpler processes: first identifying highly similar basic blocks and then aligning the instructions in each pair of similar blocks. By operating on basic blocks, we could greatly reduce the length of the sequences to be aligned and the associated compilation and memory overhead. Furthermore, by making profitability decisions for each pair of basic blocks separately, we could avoid merging unprofitable pairs. The rest of this chapter shows how we use such an approach to overcome the weaknesses of SalSSA and make function merging practical for optimizing large programs.

6.2 Hybrid Function Merging

In this section, we propose HyFM (Hybrid Function Merging), a novel function merging technique that can operate on all functions regardless of their size with little to no compilation overheads. To achieve this goal, we rely on the insights discussed in Section 6.1. Our solution is three-fold: 1) We introduce an alignment strategy that works on the level of basic blocks, without crossing their boundaries, leading to faster and less memory demanding alignment; 2) We incorporate a multi-tier profitability analysis that allows us to bail out from unprofitable merging attempts even before code generation; 3) We introduce a linear pairwise alignment for basic blocks of the same size that produces good results on highly similar blocks. This technique can be enabled as an alternative to the quadratic sequence alignment algorithm. Both techniques have their place, offering different trade-offs.

6.2.1 Overview

For all candidate functions and basic blocks, we generate a fixed-vector representation, namely, their fingerprint [49, 50]. We match each function with its most similar avail-

able function, the one with the shortest fingerprint distance. Instead of aligning their linearized representations directly, we work at the basic block level. We pair similar basic blocks of the two functions based on their fingerprint distances. We align the instructions in these paired basic blocks using either the Needleman-Wunsch alignment [44] or our linear pairwise alignment strategy. We employ the first-tier profitability analysis on each alignment. If the cost model deems it unprofitable, we skip the pair. The pairing of basic blocks, the alignment, and the first-tier profitability analysis are executed in rounds, in a greedy manner. That is, the first profitable pairing is taken, however, unprofitable paired blocks are freed for another pairing, if necessary.

Once all basic blocks have been processed, we combine the block alignments into a function-wide one and we produce the merged function using the same code generation proposed for SalSSA [50]. If no profitable pair of basic blocks was found, we bail out before code generation. Finally, we perform the second-tier profitability analysis, which is the same used by SalSSA, to decide whether replacing the original functions by the merged one reduces code size. If not, we reject the merged function and we keep the original ones.

For brevity, the rest of the discussion will focus on how HyFM differs from previous approaches.

6.2.2 Pairing Similar Basic Blocks

We pair similar basic blocks based on distance of their fingerprints. This pairing process is similar to the search strategy used for pairing functions [49]. We use the same fingerprint, a fixed-size vector of integers with the frequency count of each opcode. It can be used to represent any piece of code, from basic blocks to whole functions.

The overall idea is that for each block in one function we select a block from the other function that minimizes the Manhattan distance between their fingerprints. Formally, given a block $B_1 \in F_1$, where F_1 is the set of all blocks from function one, B_1 is paired with a block $B_m \in F_2$ such that:

$$d(B_1, B_m) = \min\{d(B_1, B_2) : B_2 \in F_2\}$$

where $d(B_1, B_2)$ represents the distance between the fingerprints of the basic blocks B_1 and B_2 .

After pairing two basic blocks, B_1 and B_2 , they have their instructions aligned (see Section 6.2.3) and their merging profitability estimated (see Section 6.2.4). If

they are deemed profitable, both blocks are removed from their respective working list. Otherwise, only B_1 is removed from the working list of blocks from F_1 , i.e., B_2 can still be paired with another block, but not B_1 . In other words, basic blocks from function F_1 are paired only once, even if its alignment is deemed unprofitable. As a result, given two input functions, this pairing process is quadratic on their number of basic blocks. This number is usually much smaller than the number of instructions in the function, so the cost of pairing is much lower than the cost of aligning whole functions in SalSSA, despite both being quadratic. For very large numbers of basic blocks, efficient nearest neighbor search techniques could keep the cost low but this was not needed in our experiments.

HyFM pre-computes the fingerprint of every basic block in the input functions, which is a single linear cost over all their basic blocks and instructions. Meanwhile, the distance between two fingerprints is computed in constant time, since the number of opcodes is a small constant.

6.2.3 Aligning Paired Basic Blocks

Basic blocks already represent a linearized sequence of instructions. Any sequence alignment algorithm can be used on them the same way they can be used on linearized functions. The globally optimal Needleman-Wunsch [44] used by SalSSA remains a good choice. It may be quadratic in both space and time on the length of the sequences but basic block sequences are usually much shorter than functions, making the cost of alignment lower than in previous approaches.

Our observations in Section 6.1, though, indicate that a globally optimal algorithm might be an excessive solution. Profitable sequences tend to be highly similar, so aligning them is usually straightforward. Based on this insight, we implemented a linear alignment algorithm. Its assumption is that profitable pairs of blocks are almost identical in terms of opcodes differing only in a few individual cases. This translates into a pairwise alignment of same size basic blocks where only corresponding instructions in the two blocks can match. Figure 6.3 illustrates two examples of basic blocks aligned using our strategy. It also includes the costs estimated by our profitability analysis, which we discuss in Section 6.2.4.

Restricting alignment to basic blocks of the same size has the added benefit that it simplifies the pairing strategy. We only have to consider fingerprints for basic blocks of the same size, so we group them by block size and we restrict our search in the right

%sw.bb	%entry	0)+1
%v1 = gep %this, 0, 5	%x1 = alloca	2	
%v2 = bitcast %v1	%x2 = gep %this, 0, 1	2)+2
%v3 = load %v2	%x3 = load %x2	1	
%v4 = icmp eq %v3, 0	%x4 = icmp eq %x3, 73	1	
br %v4, Lb3, Lb2	br %x4, Lb3, Lb2	1	
Merged Cost: 10 ✓			

(a) A profitable alignment. Both *OriginalCost* and *MergedCost* are 10. The final score is $OriginalCost - MergedCost = 0$.

%entry	%entry	0)+1
%v1 = gep %this, 0, 21	%x1 = extract %pn	2	
%v2 = bitcast %v1	%x2 = call catch(%e)	2)+2
store %c, %v2	%x3 = gep %x2, 8	2	
%v3 = gep %this, 0, 0	%x4 = bitcast %x3	2)+1
%v4 = load %v3	%x5 = load %x4	1	
%x = invoke create(%x1)	invoke printfm(...)	2	
Merged Cost: 15 ✗			

(b) An unprofitable alignment. *OriginalCost* is 0 and *MergedCost* is 15. The final score is -3 . A negative score means it is unprofitable.

Figure 6.3: Two examples of the pairwise alignment. Only instructions in corresponding positions are aligned. Instructions match if they have the same opcode.

group. In the worst case, all basic blocks would have the same size and the search would remain quadratic on the number of basic blocks, as discussed in Section 6.2.2. However, this is unlikely to happen in large functions, which is where the number of basic blocks might be a problem. Overall, this solution is lean on memory usage and usually the fastest for aligning paired basic blocks, as corroborated by our evaluation in Section 6.3.

6.2.4 Multi-Tier Profitability Analysis

HyFM incorporates a multi-tier profitability analysis that enables it to bail out early from an unprofitable merge operation. The first tier consists of a simple analysis applied on each pair of basic blocks selected for alignment, either accepting or rejecting the alignment between two blocks. The second tier consists of the same profitability analysis that is also performed by FMSA and SalSSA, which is responsible for evaluating whether the merged function is smaller than the original input functions.

The last column of Figure 6.3 shows how the first tier analysis is employed alongside the pairwise alignment strategy. The same analysis can also be applied on pairs of basic blocks aligned using the Needleman-Wunsch algorithm. The analysis tries

to estimate the cost of merging, the total number of instructions that will be necessary for merging the aligned blocks. If two instructions match, then a single instruction is needed (i.e., a cost of +1 is assigned to this entry). If they mismatch, then both instructions are needed (i.e., a cost of +2 is assigned to this entry). Moreover, we need extra instructions to transition from matching subsequences to mismatching ones, and vice versa. This is represented by the arrows in Figure 6.3. One branch instruction is needed to split control flow into two mismatching instructions, while two branch instructions are needed to join it back into a matching pair of instructions. The *MergedCost* is the sum of all these costs. The profitability score is defined as $OriginalCost - MergedCost$, where *OriginalCost* is simply the number of instructions in the original basic blocks. Therefore, a negative profitability score means that merging those two basic blocks is unprofitable. When this is the case, we ignore the alignment.

By rejecting individual basic block alignments, we are able to decide early whether merging a pair of functions might be profitable. If we rejected all block alignments, then by definition there is no point in merging the functions. Previous approaches, without a first tier analysis, have to rely on the second tier exclusively which is applied after the functions are merged.

6.2.5 Independence from Code Layout

Unlike all prior techniques, HyFM is able to merge similar basic blocks regardless of their position in the control-flow graphs from the input functions. Figure 6.4 shows an example of two functions that all prior techniques fail to merge even though they are highly similar.




Figure 6.4: Example with code reordering extracted from the 450.soplex program.

Due to their rigid linearization strategy, both FMSA and SalSSA are unable to properly match the basic blocks of the if-else structure, resulting in sub-optimal merged

functions that are deemed unprofitable. Their linearization traverses the control-flow graph in a canonical manner, preventing blocks from being rearranged for a better merging [49, 50]. Meanwhile, earlier techniques fail to merge this example as they are restricted to functions with identical control-flow graphs where corresponding blocks are merged [2, 19, 39].

HyFM is able to correctly pair these basic blocks. Because the basic blocks are rearranged, the label operands of the conditional branch need to be handled in order to preserve the program semantics. For swapped label operands, HyFM simply uses the optimized operand resolution proposed by SalSSA, where an `xor` operation is applied on the condition of the branch and the function identifier [50].

6.3 Evaluation

In this section, we compare HyFM against SalSSA, presented in Chapter 5. First, we evaluate the code size reduction achieved by each technique, demonstrating that our approach is on a par with SalSSA. Then we show that HyFM reduces significantly the overhead of function merging. Combined with the speedup in later stages of the compilation pipeline due to the reduced amount of code, HyFM leads to faster end-to-end compilation than a baseline with no function merging enabled. Finally, we demonstrate how our contributions reduce the memory usage by several orders of magnitude.

6.3.1 Experimental Setup

In addition to evaluating SalSSA, we also consider four variations of our technique based on two dimensions: 1) the linear Pairwise Alignment (PA) versus the quadratic Needleman-Wunsch Alignment (NW), both on a per basic block manner and 2) using a multi-tier profitability analysis versus using only the standard profitability analysis from SalSSA, which is the analysis applied on the whole function after generating the merged function. As described in Section 6.2, the [PA] variant is, by construction, limited to merging only basic blocks of the same size. The [NW] variant can merge blocks of different sizes. The four variations are:

- [PA]: PA with the Multi-tier Profitability analysis.
- [PA,NMP]: PA with No Multi-tier Profitability.
- [NW]: NW with the Multi-tier Profitability.

- [NW,NMP]: NW with No Multi-tier Profitability.

To keep the comparison fair, we implemented HyFM for the same compiler as SalSSA, LLVM v11. We evaluated all techniques on the C/C++ programs from the SPEC CPU 2006 and the SPEC CPU 2017 benchmark suites [55]. The baseline in all cases is the LLVM build in full LTO mode without any function merging.

We target the Intel x86 architecture. All experiments were performed on a dedicated server with a quad-core Intel Xeon CPU E5-2650, 64 GiB of RAM, running Ubuntu 18.04.3 LTS. To minimize the effect of measurement noise, we repeated all compilation and runtime overhead experiments 5 times. We report the average values and their 95% confidence intervals.

We evaluate all approaches in terms of code size reduction, time overhead of function merging, end-to-end compilation time, and peak memory usage. To better examine the trade-off between code size reduction and compilation time, we also introduce and measure a new metric called *average reduction speed* which shows the efficiency of the optimization at reducing code size. This metric offers a single number that allows us to compare how different versions address the trade-off between compilation time and code size reduction.


Definition 1 (Average Reduction Speed) *For a given input program and optimization, let S and S_0 be the size of the program with and without the given optimization, respectively. $R = S_0 - S$ represents the reduction achieved by such optimization. Let T be the running time of the optimization pass. We define the average reduction speed as:*

$$ARS = \frac{R}{T}$$


6.3.2 Code Size Reduction

Figures 6.5 reports the reduction on the size of the linked object files produced by the compiler. While limiting alignment at a basic block granularity seems restrictive, its effect on code size is small. Even the worst performing variants of HyFM are still within 3 percentage points of SalSSA, while both [PA] and [NW] achieve good results that are on a par with SalSSA. [PA]’s code size reduction varies from 5 percentage points worse to over 10 points better than SalSSA. On average, it is within 1 percentage point of the reduction achieved by SalSSA. [NW] almost always achieves better code size reduction than [PA] and on average outperforms SalSSA by a small margin. Since

our primary aim is to reduce the high compile-time overheads of SalSSA a small loss of code reduction is acceptable.



(a) SPEC CPU 2006.



(b) SPEC CPU 2017.

Figure 6.5: Linked object size reduction over LLVM LTO when performing function merging with HyFM or SalSSA on SPEC CPU 2006 and 2017. On average, HyFM improves code size reduction.

These results indicate that the multi-tier profitability analysis is the single most important component of our approach. The two variants without the multi-tier profitability analysis, [PA,NMP] and [NW,NMP], are consistently worse than their counterparts that include this analysis, i.e. [PA] and [NW]. The multi-tier analysis contributes on average about 1 percentage point in code reduction for SPEC 2017 and more than 2 points for SPEC 2006. The multi-tier profitability analysis has an important impact in the quality of the merged function. While SalSSA lets unprofitable merged subse-

quences through as long as they are outweighed by profitable subsequences elsewhere in the merged function, HyFM filters such unprofitable subsequences out.

The next most important effect comes from the choice of alignment algorithm. Needleman-Wunsch is on average half a percentage point better than Pairwise Alignment for SPEC 2017 and about one percentage point better for SPEC 2006. Given that Pairwise Alignment only aligns blocks of the same size and does not try to discover optimal alignments, this difference is smaller than one would expect. It indicates that profitable pairs of basic blocks tend to be extremely similar if not identical, as discussed in Section 6.1. Still, Needleman-Wunsch results in more size reduction. When code size reduction is paramount, [NW] might be a better choice than [PA], but as we will see in Sections 6.3.4 and 6.3.6, there is still a trade-off to navigate.

We get the biggest improvement by [PA] compared to SalSSA and [NW] for `lbm`, where it reduces the program’s object file by 18.5%, almost 13 percentage points more than the competition. SalSSA is able to profitably merge two pairs of functions. On the other hand, [PA,NMP] chooses to perform a chained merge of the three largest functions in `lbm`, resulting in a significantly smaller binary. This is possible because [PA,NMP] is merging some nearly identical pairs of basic blocks of the same size. With the multi-tier profitability analysis, [PA] successfully identifies all four cases. [NW] fails to identify all of these cases, even though it is still better than SalSSA. This exposes existing limitations in the cost model used by our profitability analysis.

The two worst results for [PA] are for the `namd` benchmark in both SPEC 2006 and SPEC 2017. SalSSA achieves close to 7 percentage points more than [PA] in code size reduction. In both cases, Pairwise Alignment limits the number of successful merge operations. The variants using Needleman-Wunsch recover most of the lost reduction.

6.3.3 Speeding Up Function Merging

Figure 6.6 shows the speedup of HyFM relative to SalSSA. This considers only the time taken by the function merging pass, which include all stages discussed in Section 6.1.2. Our novel technique achieves an impressive speedup. For [PA] it is on average $5.28\times$ faster for SPEC 2017 and $4.59\times$ for SPEC 2006. Even in the worst case, it achieves a 50% speedup. In the best case, for the SPEC 2017 `gcc`, function merging under HyFM takes a total of 23.5 seconds instead of 302, which translates to almost $13\times$ less time.

All components of HyFM contribute towards this result but the multi-tier prof-




Figure 6.6: Speedup of the function merging pass in isolation relative to SalSSA. The multi-tier profitability analysis reduces the number of unprofitable merge operations leading to a significant speedup.

itability analysis has the most significant impact. The two variants with the multi-tier profitability analysis achieve on average three to four times higher speedups than their counterparts without it. To help us understand why, Figure 6.7 shows how the compilation time of each approach is distributed across its various stages. Even though the time spent on the alignment strategy becomes negligible with HyFM, the less optimal alignment often produces complex merged functions — code with an excessive amount of branches, phi-nodes, and operand selections — slowing down SSA reconstruction and code simplification. This effect is very pronounced for the `blender` benchmark, where both [PA,NMP] and [NW,NMP] are slower than SalSSA due to the added pressure on the SSA reconstruction algorithm, even though the alignment overhead is practically

zero. Similar effects can be observed in other benchmarks.




Figure 6.7: Breakdown of the relative runtime for the different stages of the function merging pass. All measurements are normalized by SalSSA’s total runtime on the corresponding benchmark. For every benchmark, we show SalSSA, [PA,NMP], [PA], [NW,NMP], and [NW], in this order.

Enabling the multi-tier profitability analysis counters this effect by focusing code generation exclusively on profitable blocks and functions. Most of the complex basic blocks HyFM generates are not profitable for the same reason it is expensive to process them. First-tier profitability filters them out. On top of that, most paired functions under either SalSSA or HyFM are unprofitable. SalSSA has to merge them anyway to determine their profitability. This is expensive and wasteful. Our approach, on the other hand, is able to estimate profitability early. Only function pairs with any chance of being profitable, that is pairs with at least one profitable pair of basic blocks, move forward to the expensive merge stage.

The linear pairwise alignment contributes to the performance improvement, too. The variants using pairwise alignment run on average 48% to 98% faster than their Needleman-Wunsch counterparts. The most pronounced case is for `lbm` where [PA] is around $3\times$ faster than [NW]. The blocks paired in `lbm` are longer than usual, so the quadratic Needleman-Wunsch spends significantly more time trying to align them than our linear pairwise algorithm. Figure 6.7 shows that the added pairing restrictions from [PA], to focus on blocks with higher similarities, also benefits later stages.

6.3.4 End-to-End Compilation Time


We have also analyzed separately the end-to-end compilation time because reducing code size through function merging has knock-on effects in later stages of the compilation pipeline. The first order effect is that reducing the number of functions tends to reduce compilation time. This is not guaranteed though, because merged functions may be more complex, potentially slowing down later compiler analyses and transformations. Moreover, the time spent merging functions may be so large that it negates any benefits from having fewer functions later in the pipeline.

Even though on a few occasions SalSSA reduces end-to-end compilation time, in general, its overhead is large enough to result in an overall compilation time slowdown, 9.5% to 4.1% for SPEC 2017 and 2006 respectively. In contrast, HyFM is so much faster that its compilation time overhead is matched or outweighed by the speedup in later stages. This reduction is marginal for SPEC 2017, but for SPEC 2006 [PA] reduces the average compilation time by 2.3% and [NW] by 1.6%. There is only a single case where [PA] results in a significant end-to-end slowdown, 10% for `blender`. Figure 6.7 shows that, although [PA] runs faster than SalSSA, both of them spent a significant amount of time ranking the function candidates, due to its large number of functions. Ranking alone in this case takes around 70 seconds.


Overall, we believe that this reduction in end-to-end compilation time is a very important result. While HyFM is still achieving code-size reduction on par with the state-of-the-art, it does not have the detrimental effects of SalSSA on the overall compilation process and can be safely applied.

6.3.5 Code Size and Compilation Time Trade-Off

While the multi-tier profitability analysis improves both code-size reduction and compilation speed, the choice of alignment algorithm introduces a trade-off. Pairwise alignment is better for speed, Needleman-Wunsch is better for code-size reduction. In terms of compilation efficiency, i.e. how much code size reduction we get for the effort we put in, the picture is clearer. In Figure 6.9, the *average reduction speed* suggests that [PA] achieves the ideal trade-off, with an average reduction speed of 115.3 KB/s, which is around $3\times$ greater than SalSSA's and 20% to 40% greater than [NW].



(a) SPEC CPU 2006.




(b) SPEC CPU 2017.

Figure 6.8: Normalized end-to-end compilation time for SPEC 2017 and SPEC 2006 relative to LLVM LTO.


6.3.6 Memory Usage

Another important aspect of function merging is peak memory usage. This is especially critical for an optimization designed for LTO. Compilation in full LTO mode is already memory hungry. Just keeping the whole program in memory can be a significant problem for large programs [30]. Maintaining additional information for every function and basic block could easily tip the compiler over the edge.

Figure 6.10 shows the peak memory usage (in log scale) needed for the alignment stage alone. For SalSSA, this represents simply the execution of the Needleman-Wunsch algorithm. For HyFM, the alignment stage represents both aligning each pair of basic blocks as well as the pairing these basic blocks. Our results show that [PA]



(a) SPEC CPU 2006.



(b) SPEC CPU 2017.

Figure 6.9: Average reduction speed on both SPEC 2006 and 2017.

is over three orders of magnitude better than SalSSA, while [NW] is more than two orders of magnitude better. In other words, while SalSSA requires on average 2.4 GB of memory, [PA] uses only around 610 KB and [NW] uses 5.6 MB.

The peak memory usage is especially noticeable on `gcc`, when merging its two largest functions, containing 90093 and 76265 instructions. Since SalSSA applies its quadratic sequence alignment algorithm on the linearized sequences of the whole input functions, it uses over 32 GB of memory when merging these two functions. Meanwhile, [NW] requires only around 5.6 MB for merging the same pair of input functions, even though it employs the same sequence alignment algorithm. This is because its peak memory usage is a quadratic function of the largest pair of *blocks* instead of the largest pair of *functions*. Although very large, these two functions are composed of


several thousands of very small basic blocks, so the memory overhead of Needleman-Wunsch is limited. Most of the memory consumed by [NW] in this case is actually needed for storing the basic block fingerprints. This aspect becomes evident when we compare the peak memory usage of [NW] with that of the [PA] for `gcc`. They have similarly low memory requirements, even though only one of them uses a quadratic alignment algorithm.

In other cases, where basic blocks are longer, pairwise alignment leads to a significantly lower peak memory usage compared to Needleman-Wunsch. For `parest`, for example, pairwise alignment reduces memory usage from 40 MB to 200 KB. Overall, [PA] needs around $6\times$ less memory. For smaller programs, [NW] might be a viable option but for larger ones being able to reduce memory usage to a minimum might be more important.


6.4 Conclusion

We have presented HyFM, a novel technique for compiler-based function merging. By operating on individual pairs of basic blocks, it eliminates most of the time and space overheads of SalSSA. Through its multi-tier profitability analysis, it allows the compiler to bail out early from unprofitable merging attempts saving additional compilation time.

We evaluate HyFM by applying it to SPEC CPU2006 and 2017 benchmark suites. Overall, HyFM has surpassed SalSSA in terms of compilation time, memory usage, as well as code size reduction. However, different variants of the proposed technique are better suited for different goals. If the code size is the utmost concern, HyFM [NW] is the winning strategy, but if we are looking for the most balanced trade-off between compilation-time overheads and code-size reduction, HyFM [PA] has shown better results.



(a) SPEC CPU 2006.



(b) SPEC CPU 2017.

Figure 6.10: Peak memory usage of SalSSA and HyFM variants for SPEC 2006 and 2017 in log scale. SalSSA has a peak memory usage several orders of magnitude hundreds higher than all other approaches. The pairwise alignment variants of HyFM need on average only a seventh of the memory needed by the Needleman-Wunsch variants.

Chapter 7

Conclusion

This thesis presents novel compiler optimisations based on function merging to reduce code size. A function merging optimisation consists in combining multiple functions with similarity into a single function, in order to reuse code, removing redundancies, and reducing code size.

7.1 Contributions

7.1.1 Merging Arbitrary Pairs of Functions

Chapter 4 describes a novel function merging optimisation, called FMSA. The proposed technique is the first capable of merging any two functions. In order to achieve that, our technique uses sequence alignment algorithms, from bioinformatics, to identify regions of similarity and difference between two functions. The sequence alignment algorithm produces this similarity information by searching for the best alignment between sequences, given a penalty system. The resulting aligned sequence is then used to generate a merged function that can replace both input functions.

Since we are able to merge arbitrary pairs of functions, we also need to identify which pairs of functions can be profitably merged, reducing code size. First, we use a profitability cost model that estimates the binary size of the input functions and the merged function. Two input functions can be profitably merged if replacing them by a merged function results in an overall smaller code. Then we integrate our function merging technique with a search strategy.

Trying to merge every pair of functions in order to identify the most profitable ones is prohibitive. Besides the quadratic search itself, merging two functions involve

a sequence alignment, which is quadratic on the size of the functions. Since an exhaustive search is infeasible, we need a better search strategy. First, we pre-compute the fingerprint, which is a fixed-size representation, of every function in the program. These fingerprints can be efficiently compared in constant time. For each function, the search strategy consists of ranking all other candidate functions based on a fingerprint similarity. Finally, the proposed optimisation will attempt to merge only the top ranked candidate functions.

Our optimisation is able to reduce code size by up to 25%, with an overall average of about 6%, when compared to a baseline without any function merging. Compared to prior techniques, which contain many limitations on which functions they can merge, our approach is able to outperform them by more than $2.4\times$ on code size reduction. Moreover, coupled with profiling information, our optimisation introduces no statistically significant impact on performance.

While representing a leap forward, experiments show that FMSA fails to reduce code size in some cases where it would be intuitively expected to work. Chapters 5 and 6 addresses major limitations from FMSA, improving both its capabilities for reducing code size as well as compilation-time overheads.

7.1.2 Effective Code Generator for the SSA Form

In order to simplify its code generator, the technique proposed in Chapter 4 replaced all phi-nodes with memory operations by first applying register demotion. This tends to almost double the function size, increasing compilation time and hindering function merging, as it expects that applying register promotion to the merged function will reverse the negative effect of the earlier register demotion. This is often not possible because function merging can add complexity to the memory operations, resulting in unnecessarily larger merged functions.

Chapter 5 describes a novel approach for merging functions which is capable of effectively handling the SSA form. The proposed approach, called SalSSA, achieves this with a new code generator. Instead of translating the aligned sequences directly into a merged function, SalSSA generates code from the input control-flow graphs, using the alignment only to specify pairs of matching labels and instructions. The generator then produces code top-down, starting with the control flow graph of the merged function, then populating with instructions, arguments and labels, and finally with phi-nodes which maintain the correct flow of data.

Chapter 5 also introduces a novel phi-node coalescing optimisation which is integrated with function merging in order to produce fewer phi-nodes and selections in the merged functions. This optimisation pairs disjoint values coming from each one of the input functions so that they can share the same phi-nodes.

7.1.3 Merging Functions One Block at a Time

Both FMSA and SalSSA has a coarse-grained strategy for merging functions. They apply the quadratic sequence alignment algorithm on the linearised input functions. The linearisation specifies an ordering for all the basic blocks from a function, producing a sequence with all the instructions of the function in that order. Different orderings of the basic blocks would result in different alignments and therefore different merged functions. Moreover, they also apply a profitability analysis once the merged functions has been completely generated, either accepting or rejecting it completely.

Chapter 6 introduces HyFM, our function merging technique that works on the basic block level. Since each basic block already represents a sequence of instructions, HyFM does not require a linearisation step. It pairs similar basic blocks using a fingerprint-based distance metric. These pairs are aligned either using the Needleman-Wunsch algorithm or a linear one that works in a pairwise manner.

Chapter 6 also introduces a multi-tier profitability analysis. Before generating the merged function, it uses a cost model to analyse the alignment for each pair of basic blocks. Alignments deemed unprofitable are discarded. This fine-grained tier of the profitability analysis improves the quality of the merged functions produced by HyFM, but also enables it to bail out early from unprofitable merging attempts.

Overall, HyFM produces similar code size reductions to SalSSA, while running significantly faster. Unlike our prior techniques, HyFM reduces end-to-end compilation time.


7.2 Future Work

For future work, we plan to focus on improving the ranking mechanism to reduce compilation time. In order to avoid code size degradation, we also plan to improve the compiler's built-in static cost model for code size estimation. We also plan to work on the linearisation of the candidate functions, allowing instruction reordering to maximise the number of matches between the functions. Finally, we also plan to

incorporate instruction reordering into function merging to maximise the number of matches between the functions regardless of the original code layout. It would also be interesting to investigate the application of phi-node coalescing outside function merging. We envisage further improvements can be achieved by integrating the function-merging optimisation into a summary-based link-time optimisation framework, such as ThinLTO [30] in LLVM. As a future work, we can also analyze the interaction between function merging and other optimisations such as inlining, outlining, and code splitting.

7.2.1 Handling Code Reordering at the Instruction Level

All existing function merging techniques rely on the order in which instructions appear inside the basic blocks and their arrangement, failing to profitably merge equivalent functions when confronted even with the smallest variations on the ordering of instructions and basic blocks. Figure 7.1a shows an example of two functions that all existing techniques fail to merge even though they are obviously identical. Our current technique produces the merged function shown in Figure 7.1b, which is considered unprofitable as it is unable to reduce code size. Note that all indexing computation is duplicated, including the increment operation, resulting in a merged function that is unnecessarily bigger than the two original functions together. We need more powerful graph, rather than sequence, alignment techniques to better identify and merge differently ordered but semantically equivalent code.



```

int foo(int *V, int i){      int bar(int *V, int i){
    int t0 = V[i];           int t1 = V[i+1];
    int t1 = V[i+1];         int t0 = V[i];
    return (t0+t1);          return (t0+t1);
}                             }

```

(a) Two equivalent functions.

```

int m_foo_bar(int fid, int *V, int i) {
    int idx0 = fid?i:(i+1);
    int var0 = V[idx0];
    int idx1 = fid?(i+1):i;
    int var1 = V[idx1];
    return (var0+var1);
}

```

(b) Merged function currently produced by FMSA.

Figure 7.1: Example of how even trivial reordering is poorly handled by the existing solutions.

7.2.2 Merging Across All Scopes

Existing techniques are limited to one particular scope. While function merging is applied only to whole functions, function outlining is commonly applied at the basic block level. However, equivalent code can be found within or across functions, which themselves may reside in the same source file or be spread across multiple source files. Therefore, we need to develop a novel unified optimisation capable of merging semantically equivalent code that can span anything between a single basic block up to a whole function. This unification has the extra benefit of also addressing the phase ordering problem by coordinating the merge operations on different scopes.

7.2.3 Scaling for Large Programs

Although our optimisation achieves very good results in terms of code compression, it is still unable to handle large programs in a real scenario. Its time complexity and memory usage requirements would prevent it from optimizing large programs such as web browsers, Clang/LLVM, and operating systems, as these programs tend to have many functions with several thousands of instructions. Link-time optimisation (LTO) makes this matter even worse by optimizing the code after the whole program has been linked into a single module, imposing a huge pressure on memory usage and compilation time.

The optimisation of different translation units can be distributed across different machines and merge operations locally performed in parallel. An important challenge that needs to be addressed concerns ranking and merging functions that reside in different translation units. However, parallel LTO is essential to enable its use on real programs while keeping the memory usage and compilation time acceptable. We believe these improvements can be achieved by integrating the function-merging optimisation with a summary-based link-time optimisation framework, such as ThinLTO in LLVM.

7.2.4 Powered by Deep Learning

One can investigate the use of deep learning to align two functions and better identifying what can be merged. Sequence alignment focusses only on maximizing the number of merged instructions, without necessarily minimizing the number of operand selections or branches. A smarter approach that understands how instructions interact with each other would be very beneficial.

7.2.5 Avoiding Performance Overheads

For many real applications, it is desirable to achieve a good balance between code size and performance. Preliminary results show that performance degradation can be completely avoided by using profiling to guide the merging decisions. One can avoid adding branches inside hot execution paths, therefore avoiding performance penalties. Although hot code can be merged, it is important to minimise unnecessary branches when merging hot code. It would be interesting to develop a profile-guided optimisation that automatically identifies the best trade-off between code-size and performance.

7.2.6 Less Memory Usage by JIT

Ahead-of-time and just-in-time (JIT) compilation have completely different requirements. Function merging can be used to reduce the amount of memory used by programs running on a JIT environment. However, our solution must be adapted to the requirements that are specific to a JIT environment, as JIT compilers have the extra challenge of having to optimise the code as fast as possible. This would require the development of completely new algorithms for ranking and aligning functions that are suitable for this application domain. Another possibility is to exploit the fact multiple programs may be simultaneously running on the same JIT environment and merge code across different programs, reducing the overall memory usage.

Bibliography

- [1] Microsoft Visual Studio. Identical COMDAT folding. <https://msdn.microsoft.com/en-us/library/bxwfs976.aspx>, 2020.
- [2] The LLVM Compiler Infrastructure. MergeFunctions pass, how it works. <http://llvm.org/docs/MergeFunctions.html>, 2020.
- [3] T. M. Ahmed, W. Shang, and A. E. Hassan. An empirical study of the copy and paste behavior during development. In *2015 IEEE/ACM 12th Working Conference on Mining Software Repositories*, pages 99–110, 2015.
- [4] Rafael Auler, Carlos Eduardo Millani, Alexandre Brisighello, Alisson Linhares, and Edson Borin. Handling IoT platform heterogeneity with COISA, a compact OpenISA virtual platform. *Concurrency and Computation: Practice and Experience*, 29(22):e3932, 2017.
- [5] Preston Briggs and Keith D. Cooper. Effective partial redundancy elimination. In *Proceedings of the ACM SIGPLAN 1994 Conference on Programming Language Design and Implementation, PLDI '94*, pages 159–170, New York, NY, USA, 1994. ACM.
- [6] Preston Briggs, Keith D. Cooper, and L. Taylor Simpson. Value numbering. *Software: Practice and Experience*, 27(6):701–724, 1997.
- [7] Humberto Carrillo and David Lipman. The multiple sequence alignment problem in biology. *SIAM J. Appl. Math.*, 48(5):1073–1082, October 1988.
- [8] Milind Chabbi, Jin Lin, and Raj Barik. An experience with code-size optimization for production iOS mobile applications. In *IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*, pages 1–12, US, 2021. IEEE Press.

- [9] John Cocke. Global common subexpression elimination. In *Proceedings of a Symposium on Compiler Optimization*, pages 20–24, New York, NY, USA, 1970. ACM.
- [10] B. Coutinho, D. Sampaio, F. M. Q. Pereira, and W. Meira Jr. Divergence analysis and optimizations. In *2011 International Conference on Parallel Architectures and Compilation Techniques*, pages 320–329, Oct 2011.
- [11] R. Cytron, J. Ferrante, B. K. Rosen, M. N. Wegman, and F. K. Zadeck. An efficient method of computing static single assignment form. In *Proceedings of the 16th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL '89, pages 25–35, New York, NY, USA, 1989. ACM.
- [12] Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13(4):451–490, October 1991.
- [13] Krzysztof Czarnecki. *Generative programming-principles and techniques of software engineering based on automated configuration and fragment-based component models*. PhD thesis, Verlag nicht ermittelbar, 1999.
- [14] Christophe de Dinechin. C++ exception handling. *IEEE Concurrency*, 8(4):72–79, October 2000.
- [15] Saumya K. Debray, William Evans, Robert Muth, and Bjorn De Sutter. Compiler techniques for code compaction. *ACM Trans. Program. Lang. Syst.*, 22(2):378–415, March 2000.
- [16] Dirk Draheim, Christof Lutteroth, and Gerald Weber. An analytical comparison of generative programming technologies. 2004.
- [17] A. Dreweke, M. Worlein, I. Fischer, D. Schell, T. Meinl, and M. Philippsen. Graph-based procedural abstraction. In *International Symposium on Code Generation and Optimization (CGO'07)*, pages 259–270, March 2007.
- [18] Tobias J.K. Edler von Koch, Igor Böhm, and Björn Franke. Integrated instruction selection and register allocation for compact code generation exploiting

- freeform mixing of 16- and 32-bit instructions. In *Proceedings of the 8th Annual IEEE/ACM International Symposium on Code Generation and Optimization*, CGO '10, pages 180–189, New York, NY, USA, 2010. ACM.
- [19] Tobias J.K. Edler von Koch, Björn Franke, Pranav Bhandarkar, and Anshuman Dasgupta. Exploiting function similarity for code size reduction. In *Proceedings of the 2014 SIGPLAN/SIGBED Conference on Languages, Compilers and Tools for Embedded Systems*, LCTES '14, pages 85–94, New York, NY, USA, 2014. ACM.
- [20] CD Edwards Jr, MJ Amato, JD Baker, NJ Barba, J Balaram, EJ Brandon, P Briggs, BD Davis, JC Day, A Freeman, et al. Emerging capabilities for mars exploration. *Planetary Science Decadal*, 2020.
- [21] Sebastiana Etzo and Guy Collender. The mobile phone ‘revolution’ in Africa: Rhetoric or reality? *African Affairs*, 109(437):659–668, 2010.
- [22] L. Ghica and N. Tapus. Optimized retargetable compiler for embedded processors - gcc vs llvm. In *2015 IEEE International Conference on Intelligent Computer Communication and Processing (ICCP)*, pages 103–108, 2015.
- [23] M. R. Guthaus, J. S. Ringenberg, D. Ernst, T. M. Austin, T. Mudge, and R. B. Brown. MiBench: A free, commercially representative embedded benchmark suite. In *Proceedings of the Fourth Annual IEEE International Workshop on Workload Characterization. WWC-4 (Cat. No.01EX538)*, pages 3–14, Dec 2001.
- [24] Andreas Haas, Andreas Rossberg, Derek L. Schuff, Ben L. Titzer, Michael Holman, Dan Gohman, Luke Wagner, Alon Zakai, and JF Bastien. Bringing the web up to speed with webassembly. In *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI 2017, page 185–200, New York, NY, USA, 2017. Association for Computing Machinery.
- [25] Stuart L. Hart and Clayton M. Christensen. The great leap: Driving innovation from the base of the pyramid. *MIT Sloan Management Review*, 44(1):51–56, Fall 2002.
- [26] Glenn Hickey and Mathieu Blanchette. A probabilistic model for sequence alignment with context-sensitive indels. In *Proceedings of the 15th Annual In-*

- ternational Conference on Research in Computational Molecular Biology*, RE-COMB'11, pages 85–103, Berlin, Heidelberg, 2011. Springer-Verlag.
- [27] Desmond G Higgins and Paul M Sharp. Fast and sensitive multiple sequence alignments on a microcomputer. *Bioinformatics*, 5(2):151–153, 1989.
- [28] P. Jablonski and D. Hou. Aiding software maintenance with copy-and-paste clone-awareness. In *2010 IEEE 18th International Conference on Program Comprehension*, pages 170–179, 2010.
- [29] Ciera Jaspan, Matthew Jorde, Andrea Knight, Caitlin Sadowski, Edward K. Smith, Collin Winter, and Emerson Murphy-Hill. Advantages and disadvantages of a monolithic repository: A case study at google. In *Proceedings of the 40th International Conference on Software Engineering: Software Engineering in Practice*, ICSE-SEIP '18, page 225–234, New York, NY, USA, 2018. Association for Computing Machinery.
- [30] Teresa Johnson, Mehdi Amini, and Xinliang David Li. Thinlto: Scalable and incremental lto. In *Proceedings of the 2017 International Symposium on Code Generation and Optimization*, CGO '17, page 111–121. IEEE Press, 2017.
- [31] S. L. Keoh, S. S. Kumar, and H. Tschofenig. Securing the internet of things: A standardization perspective. *IEEE Internet of Things Journal*, 1(3):265–275, June 2014.
- [32] Miryung Kim, L. Bergman, T. Lau, and D. Notkin. An ethnographic study of copy and paste programming practices in oopl. In *Proceedings. 2004 International Symposium on Empirical Software Engineering, 2004. ISESE '04.*, pages 83–92, 2004.
- [33] Jens Knoop, Oliver Rüthing, and Bernhard Steffen. Partial dead code elimination. In *Proceedings of the ACM SIGPLAN 1994 Conference on Programming Language Design and Implementation*, PLDI '94, pages 147–158, New York, NY, USA, 1994. ACM.
- [34] J. Kolek, Z. Jovanović, N. Šljivić, and D. Narančić. Adding micromips backend to the llvm compiler infrastructure. In *2013 21st Telecommunications Forum Telfor (TELFOR)*, pages 1015–1018, 2013.

- [35] Doug Kwan, Jing Yu, and B. Janakiraman. Google's C/C++ toolchain for smart handheld devices. In *Proceedings of Technical Program of 2012 VLSI Technology, System and Application*, pages 1–4, April 2012.
- [36] C. Lattner and V. Adve. LLVM: A compilation framework for lifelong program analysis transformation. In *Proceedings of the International Symposium on Code Generation and Optimization (CGO)*, pages 75–86, 2004.
- [37] Chris Lattner, Jacques Pienaar, Mehdi Amini, Uday Bondhugula, River Riddle, Albert Cohen, Tatiana Shpeisman, Andy Davis, Nicolas Vasilache, and Oleksandr Zinenko. Mlir: A compiler infrastructure for the end of moore's law. *arXiv preprint*, 2020.
- [38] Rahman Lavaee, John Criswell, and Chen Ding. Codestitcher: Inter-procedural basic block layout optimization. In *Proceedings of the 28th International Conference on Compiler Construction*, CC 2019, page 65–75, New York, NY, USA, 2019. Association for Computing Machinery.
- [39] Martin Liška. Optimizing large applications. *arXiv preprint arXiv:1403.6997*, 2014.
- [40] Gábor Lóki, Ákos Kiss, Judit Jász, and Árpád Beszédes. Code factoring in GCC. In *Proceedings of the 2004 GCC Developers' Summit*, pages 79–84, 2004.
- [41] Webb Miller and Eugene W. Myers. A file comparison program. *Software: Practice and Experience*, 15(11):1025–1040, 1985.
- [42] Steven S. Muchnick. *Advanced Compiler Design and Implementation*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 1998.
- [43] I. Neamtiu and T. Dumitraş. Cloud software upgrades: Challenges and opportunities. In *2011 International Workshop on the Maintenance and Evolution of Service-Oriented and Cloud-Based Systems*, pages 1–10, 2011.
- [44] Saul B. Needleman and Christian D. Wunsch. A general method applicable to the search for similarities in the amino acid sequence of two proteins. *Journal of Molecular Biology*, 48(3):443 – 453, 1970.
- [45] Pedro Plaza, Elio Sancristobal, German Carro, Manuel Castro, and Elena Ruiz. Wireless development boards to connect the world. In Michael E. Auer and

- Danilo G. Zutin, editors, *Online Engineering & Internet of Things*, pages 19–27, Cham, 2018. Springer International Publishing.
- [46] A. Pohl, B. Cosenza, and B. Juurlink. Cost modelling for vectorization on ARM. In *2018 IEEE International Conference on Cluster Computing (CLUSTER)*, pages 644–645, Sept 2018.
 - [47] Vasileios Porpodas, Rodrigo C. O. Rocha, Evgueni Brevnov, Luís F. W. Góes, and Timothy Mattson. Super-Node SLP: Optimized vectorization for code sequences containing operators and their inverse elements. In *Proceedings of the 2019 IEEE/ACM International Symposium on Code Generation and Optimization*, CGO 2019, pages 206–216, Piscataway, NJ, USA, 2019. IEEE Press.
 - [48] Vasileios Porpodas, Rodrigo C. O. Rocha, and Luís F. W. Góes. Look-ahead SLP: Auto-vectorization in the presence of commutative operations. In *Proceedings of the 2018 International Symposium on Code Generation and Optimization*, CGO 2018, pages 163–174, New York, NY, USA, 2018. ACM.
 - [49] Rodrigo C. O. Rocha, Pavlos Petoumenos, Zheng Wang, Murray Cole, and Hugh Leather. Function merging by sequence alignment. In *Proceedings of the 2019 IEEE/ACM International Symposium on Code Generation and Optimization*, CGO 2019, pages 149–163, Piscataway, NJ, USA, 2019. IEEE Press.
 - [50] Rodrigo C. O. Rocha, Pavlos Petoumenos, Zheng Wang, Murray Cole, and Hugh Leather. Effective function merging in the ssa form. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI 2020, page 854–868, New York, NY, USA, 2020. Association for Computing Machinery.
 - [51] H. Sajnani, V. Saini, J. Svajlenko, C. K. Roy, and C. V. Lopes. Sourcerercc: Scaling code clone detection to big-code. In *2016 IEEE/ACM 38th International Conference on Software Engineering (ICSE)*, pages 1157–1168, May 2016.
 - [52] Ulrik Pagh Schultz, Kim Burggaard, Flemming Gram Christensen, and Jørgen Lindskov Knudsen. Compiling Java for low-end embedded systems. In *Proceedings of the 2003 ACM SIGPLAN Conference on Language, Compiler, and Tool for Embedded Systems*, LCTES '03, pages 42–50, New York, NY, USA, 2003. ACM.

- [53] A. Sehgal, V. Perelman, S. Kuryla, and J. Schonwalder. Management of resource constrained devices in the internet of things. *IEEE Communications Magazine*, 50(12):144–149, December 2012.
- [54] T.F. Smith and M.S. Waterman. Identification of common molecular subsequences. *Journal of Molecular Biology*, 147(1):195 – 197, 1981.
- [55] SPEC. Standard Performance Evaluation Corp Benchmarks. <http://www.spec.org>, 2014.
- [56] Sriraman Tallam, Cary Coutant, Ian Lance Taylor, Xinliang David Li, and Chris Demetriou. Safe ICF: Pointer safe and unwinding aware identical code folding in gold. In *GCC Developers Summit*, 2010.
- [57] Hans Tiggeler, Tanya Vladimirova, Daixun Zheng, and Jiri Gaisler. Experiences designing a system-on-a-chip for small satellite data processing and control. In *Proceedings of International Conference on Military and Aerospace*, 2000.
- [58] M. Toomim, A. Begel, and S. L. Graham. Managing duplicated code with linked editing. In *2004 IEEE Symposium on Visual Languages - Human Centric Computing*, pages 173–180, Sept 2004.
- [59] Linda Torczon and Keith Cooper. *Engineering A Compiler*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2nd edition, 2007.
- [60] V. Porpodas, R. C. O. Rocha, and L. F. W. Góes. VW-SLP: Auto-vectorization with adaptive vector width. In *Proceedings of the 27th International Conference on Parallel Architectures and Compilation Techniques, PACT '18*, pages 12:1–12:15, New York, NY, USA, 2018. ACM.
- [61] A. Varma and S. S. Bhattacharyya. Java-through-C compilation: an enabling technology for Java in embedded systems. In *Proceedings Design, Automation and Test in Europe Conference and Exhibition*, volume 3, pages 161–166 Vol.3, Feb 2004.
- [62] V. M. Weaver and S. A. McKee. Code density concerns for new architectures. In *2009 IEEE International Conference on Computer Design*, pages 459–464, Oct 2009.

- [63] G. Yan, J. Ma, Y. Han, and X. Li. EcoUp: Towards economical datacenter upgrading. *IEEE Transactions on Parallel and Distributed Systems*, 27(7):1968–1981, 2016.
- [64] Wu Yang. Identifying syntactic differences between two programs. *Software: Practice and Experience*, 21(7):739–755, 1991.